

AutoGR: Automated Geo-Replication with Fast System Performance and Preserved Application Semantics

Jiawei Wang¹, Cheng Li¹, Kai Ma¹, Jingze Huo¹, Feng Yan², Xinyu Feng³, Yinlong Xu¹



Background

淘宝网
Taobao.com

facebook

Google

Ctrip
携程

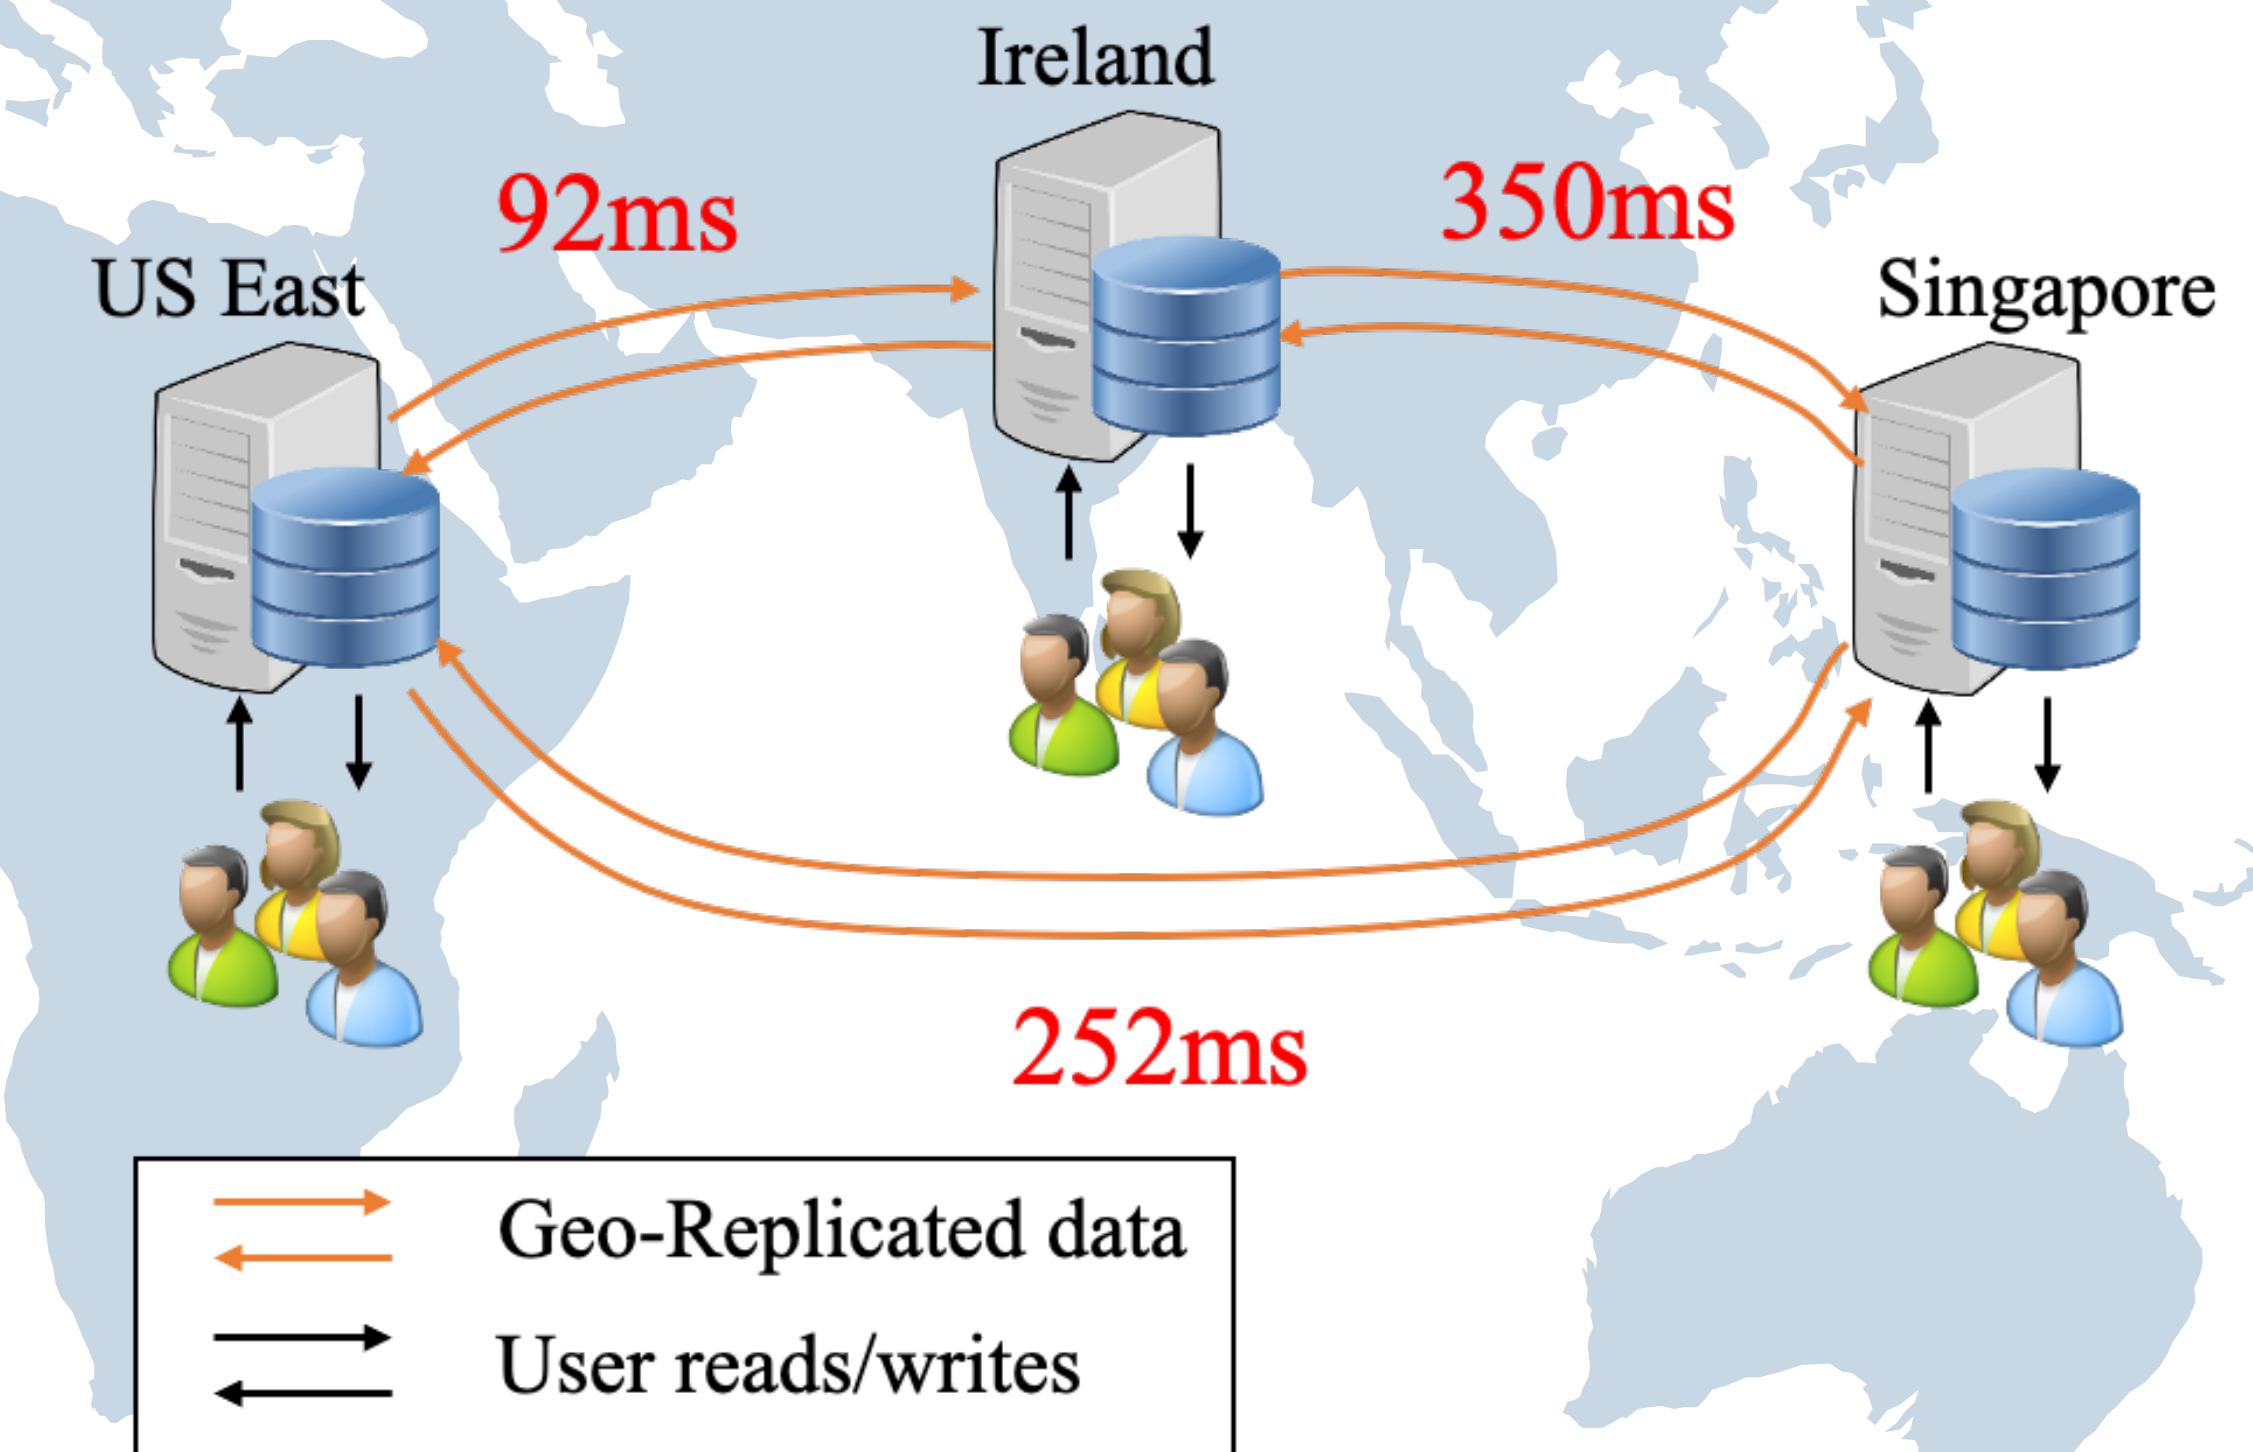
Baidu 百度

WeChat

amazon

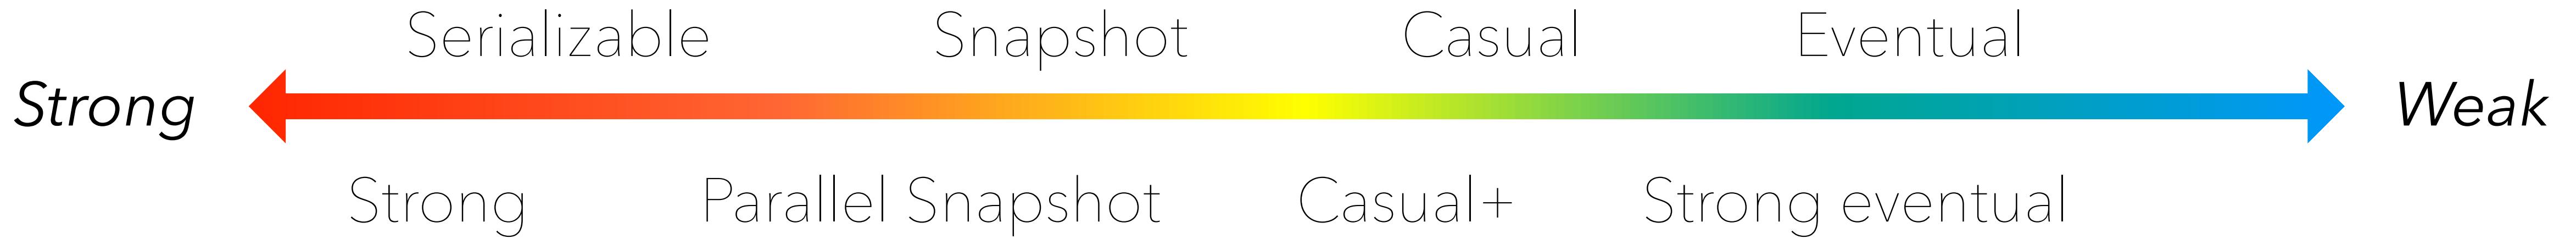
Motivation

- Geo-Replication as a major solution to cope with ever-growing user base for Internet services.

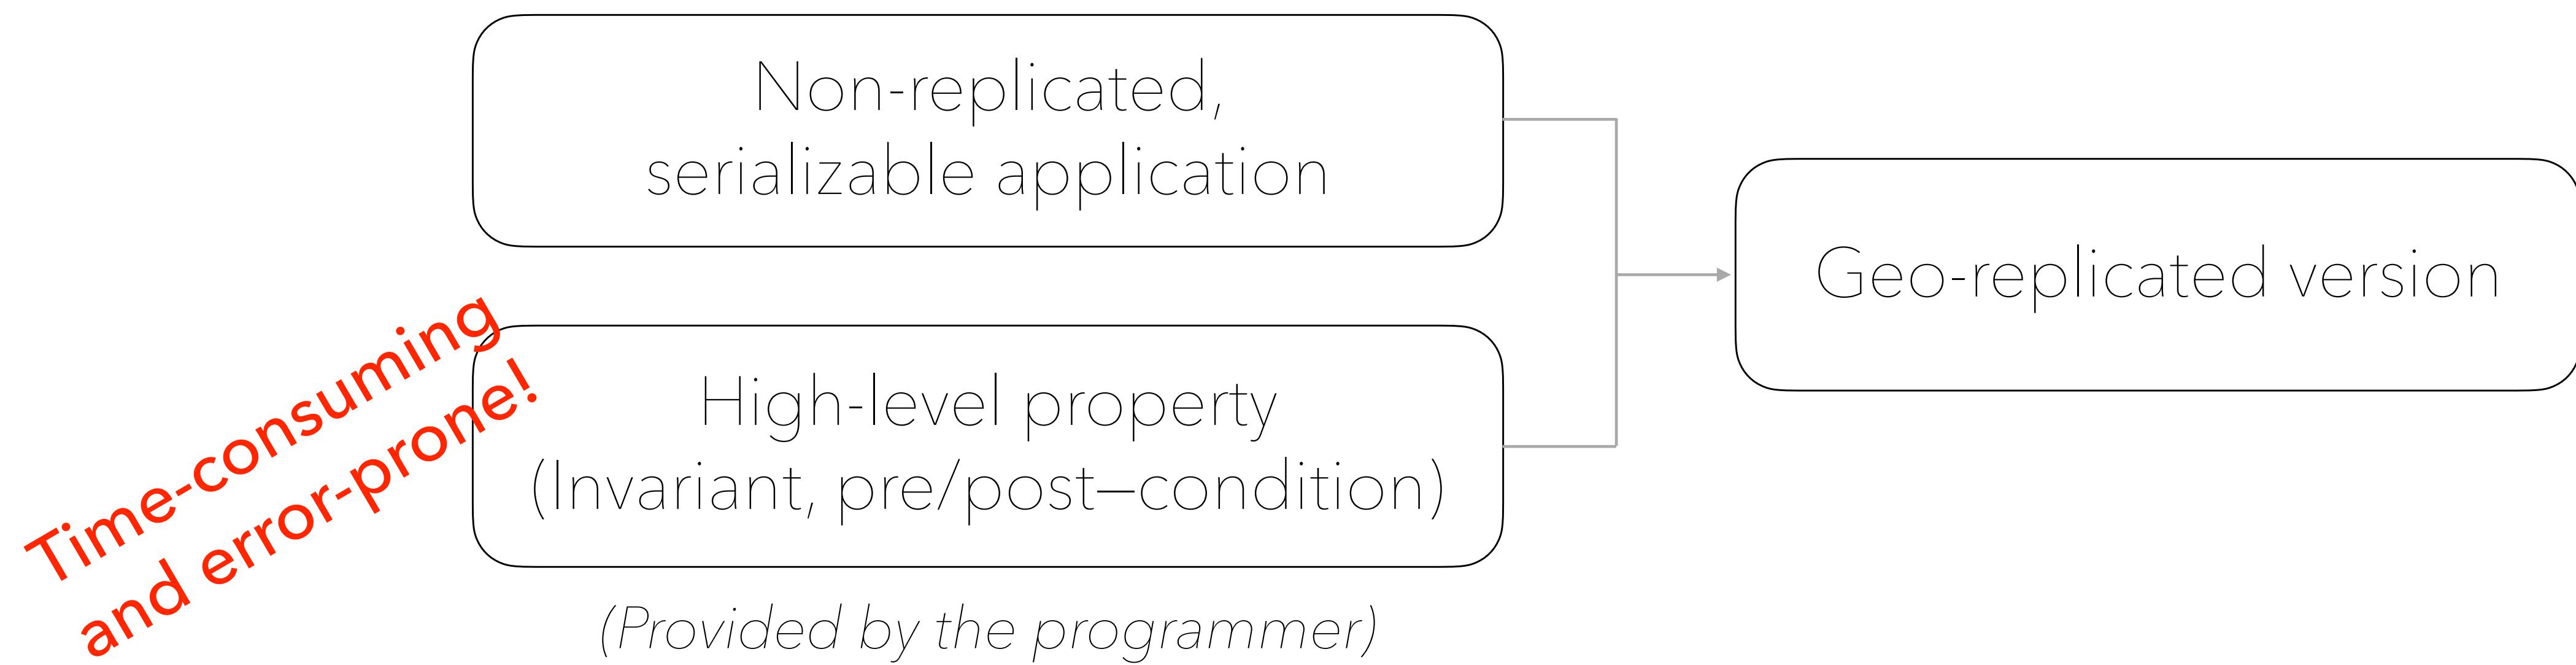


*Strong consistency introduces **high delay** to user responses.*

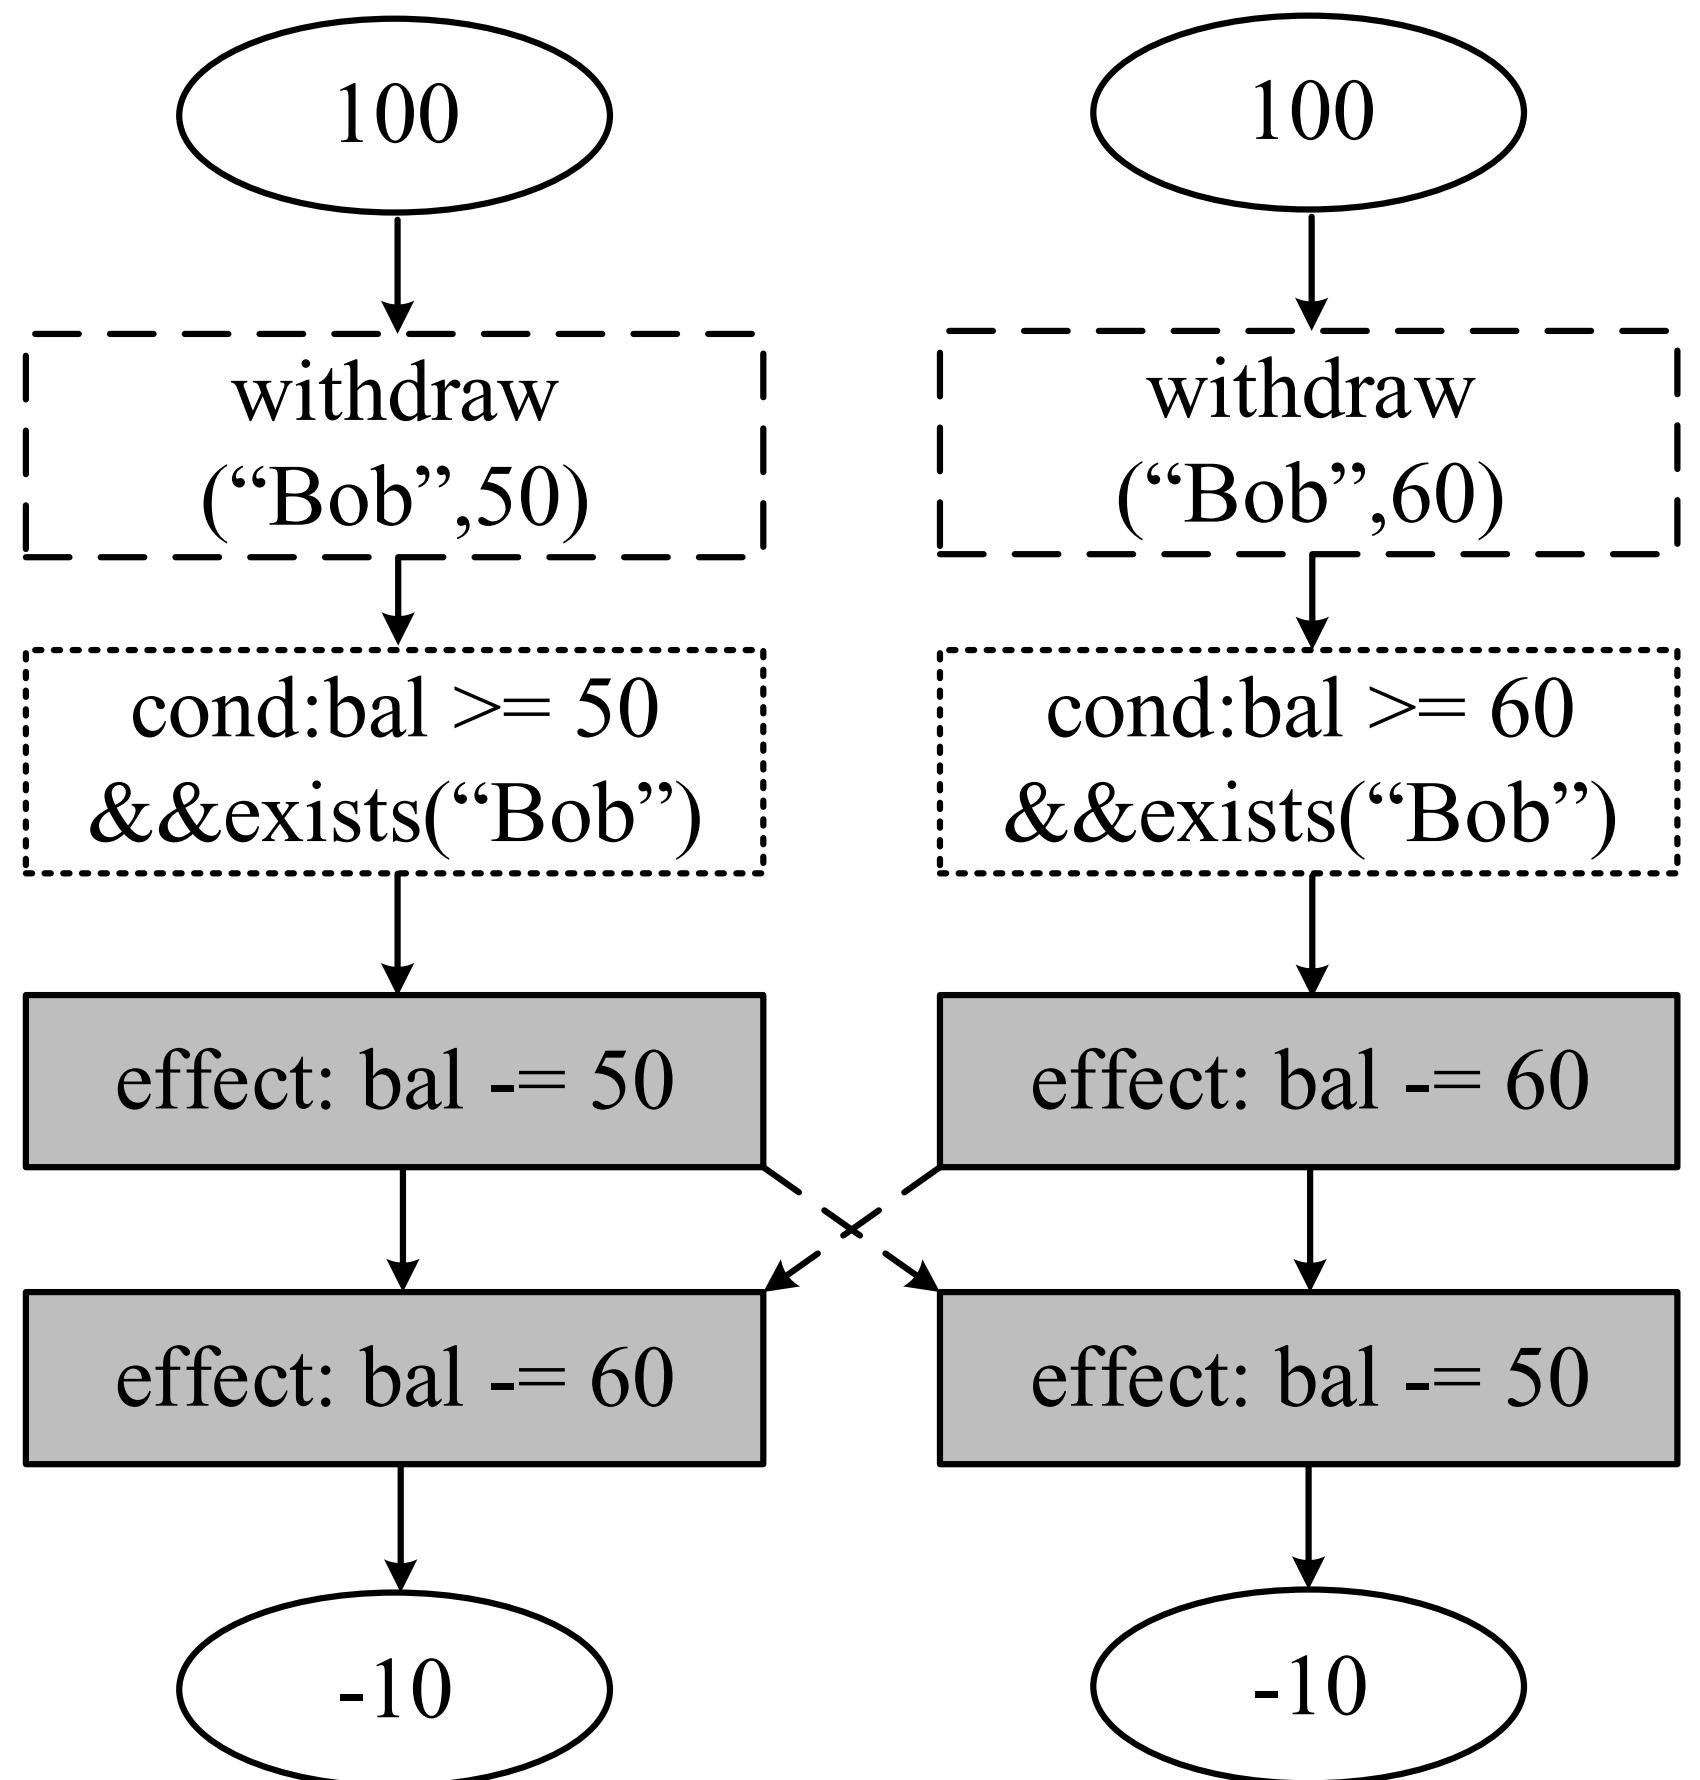
Related Work



Restriction-based fine-grained consistency model



Observation



```
void withdraw(Connection conn, String custName,  
             double amount) throws Exception {  
    PreparedStatement stmt = conn.prepareStatement(  
        "SELECT * FROM ACCOUNTS WHERE name = ?");  
    stmt.setString(1, custName);  
    ResultSet rs = stmt.executeQuery();  
    if (rs.next() == false) throw new Exception("Invalid account");  
    long custId = rs.getLong(1);  
    stmt = conn.prepareStatement(  
        "SELECT bal FROM SAVINGS WHERE custid = ?");  
    stmt.setLong(1, custId);  
    rs = stmt.executeQuery();  
    if (rs.next() == false) throw new Exception("No saving account");  
    double balance = rs.getDouble(1) - amount;  
    if (balance < 0) throw new Exception("Negative balance");  
    stmt = conn.prepareStatement(  
        "UPDATE SAVINGS SET bal =? WHERE custid=?");  
    stmt.setDouble(1, balance); stmt.setLong(2, custId);  
    stmt.executeUpdate();  
    conn.commit();  
}
```

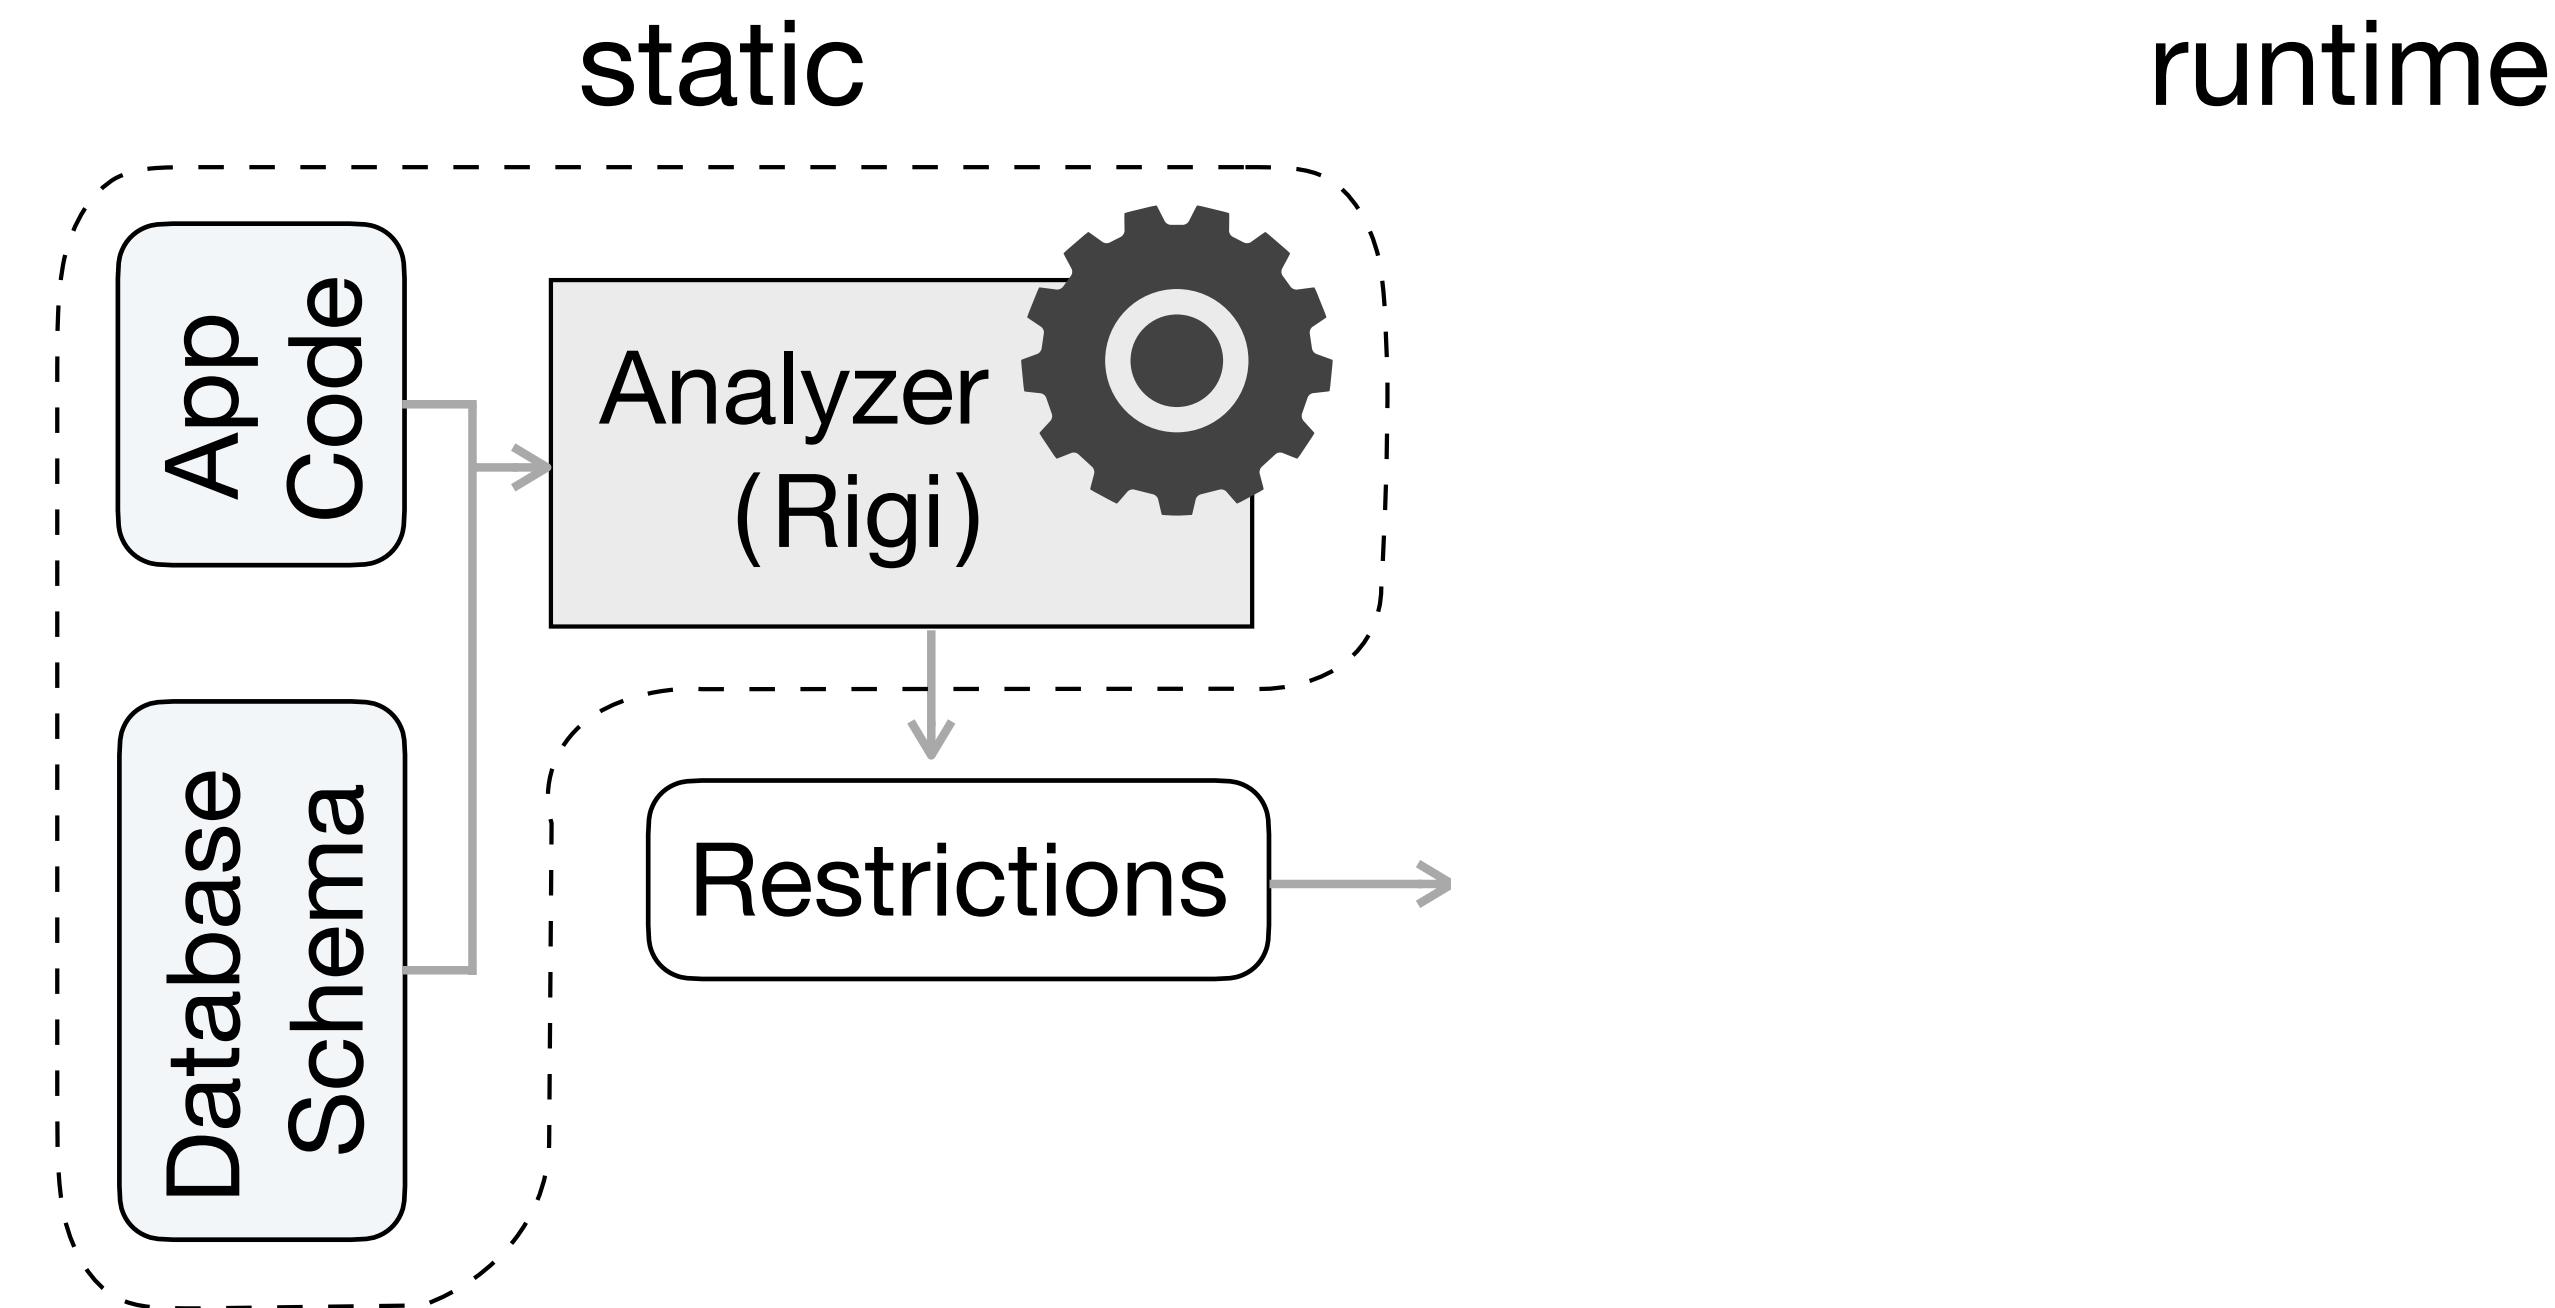
Application-specific invariants are already implicitly reflected in the programs.

AutoGR

static

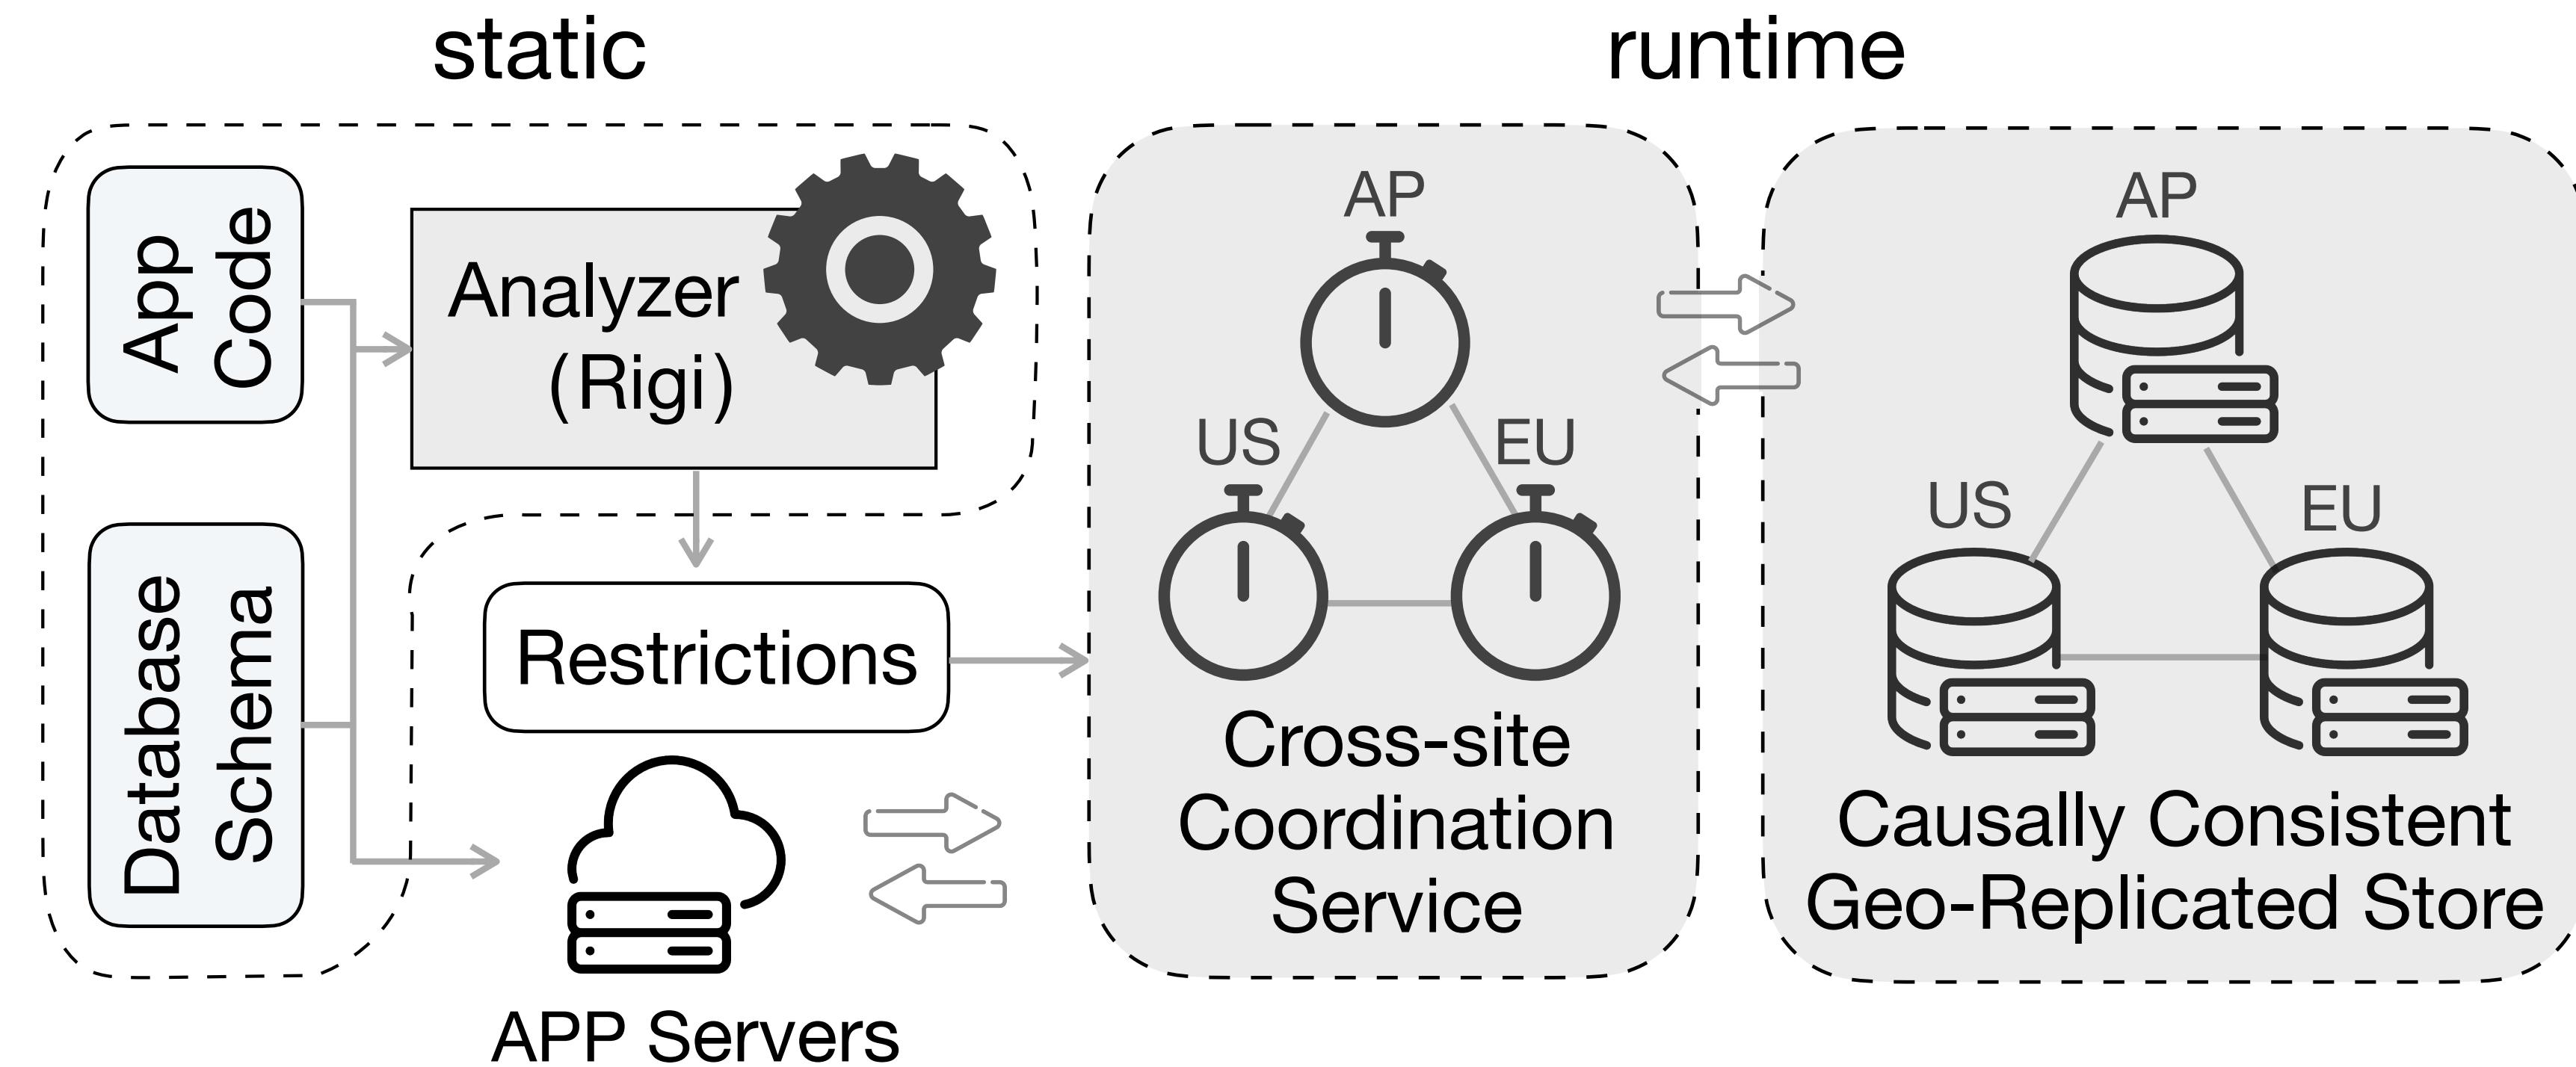
runtime

AutoGR



*The static analyzer **Rigi** identifies a minimal set of ordering restrictions that must be ensured so that the intended semantics are not violated.*

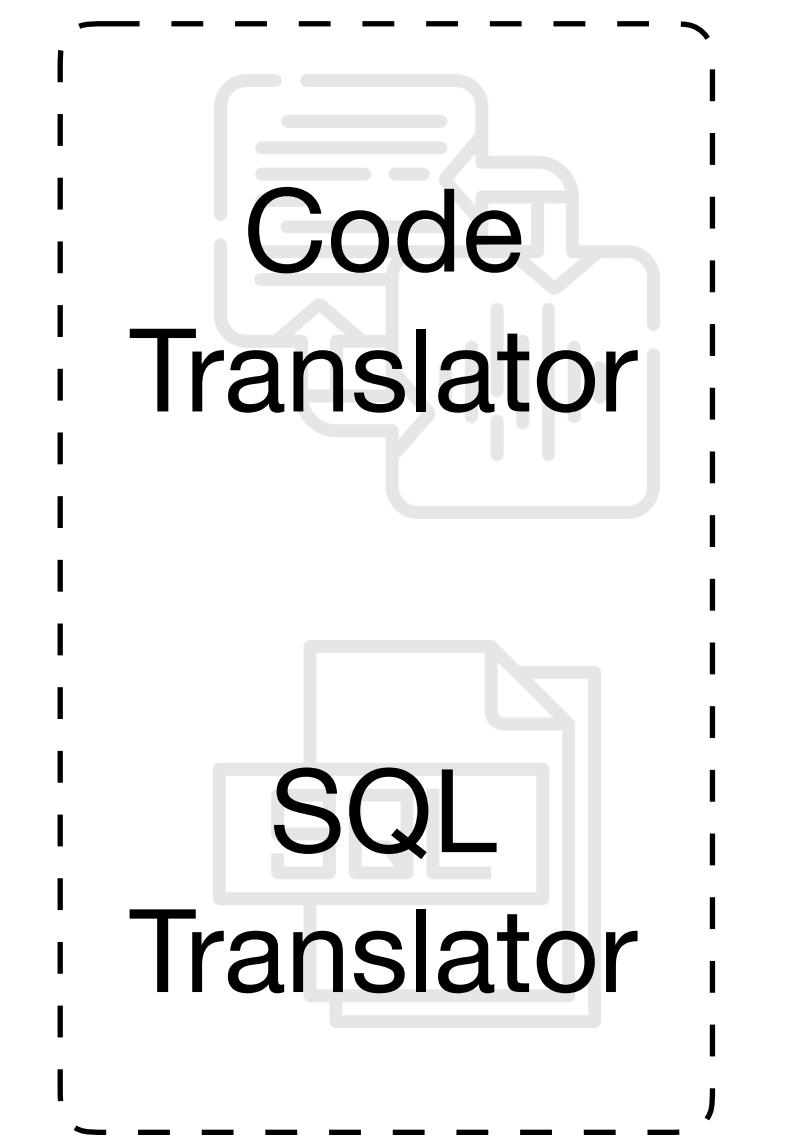
AutoGR



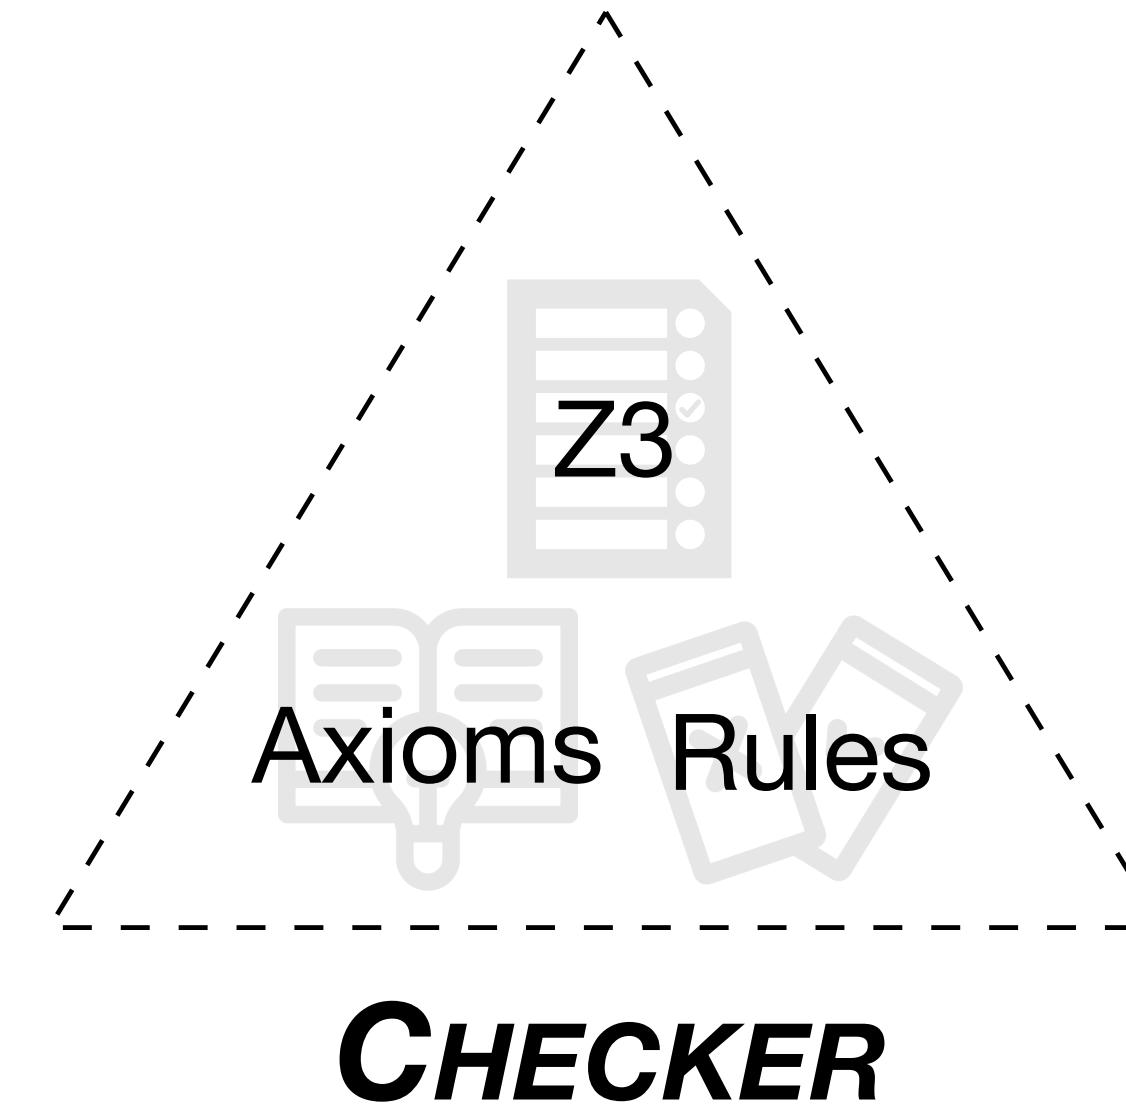
The static analyzer *Rigi* identifies a minimal set of ordering restrictions that must be ensured so that the intended semantics are not violated.

AutoGR leverages on an existing geo-replication framework *Olisipo* that enables fine-grained coordination over pairs of operations that produce conflicting side effects.

AutoGR --- Rigi

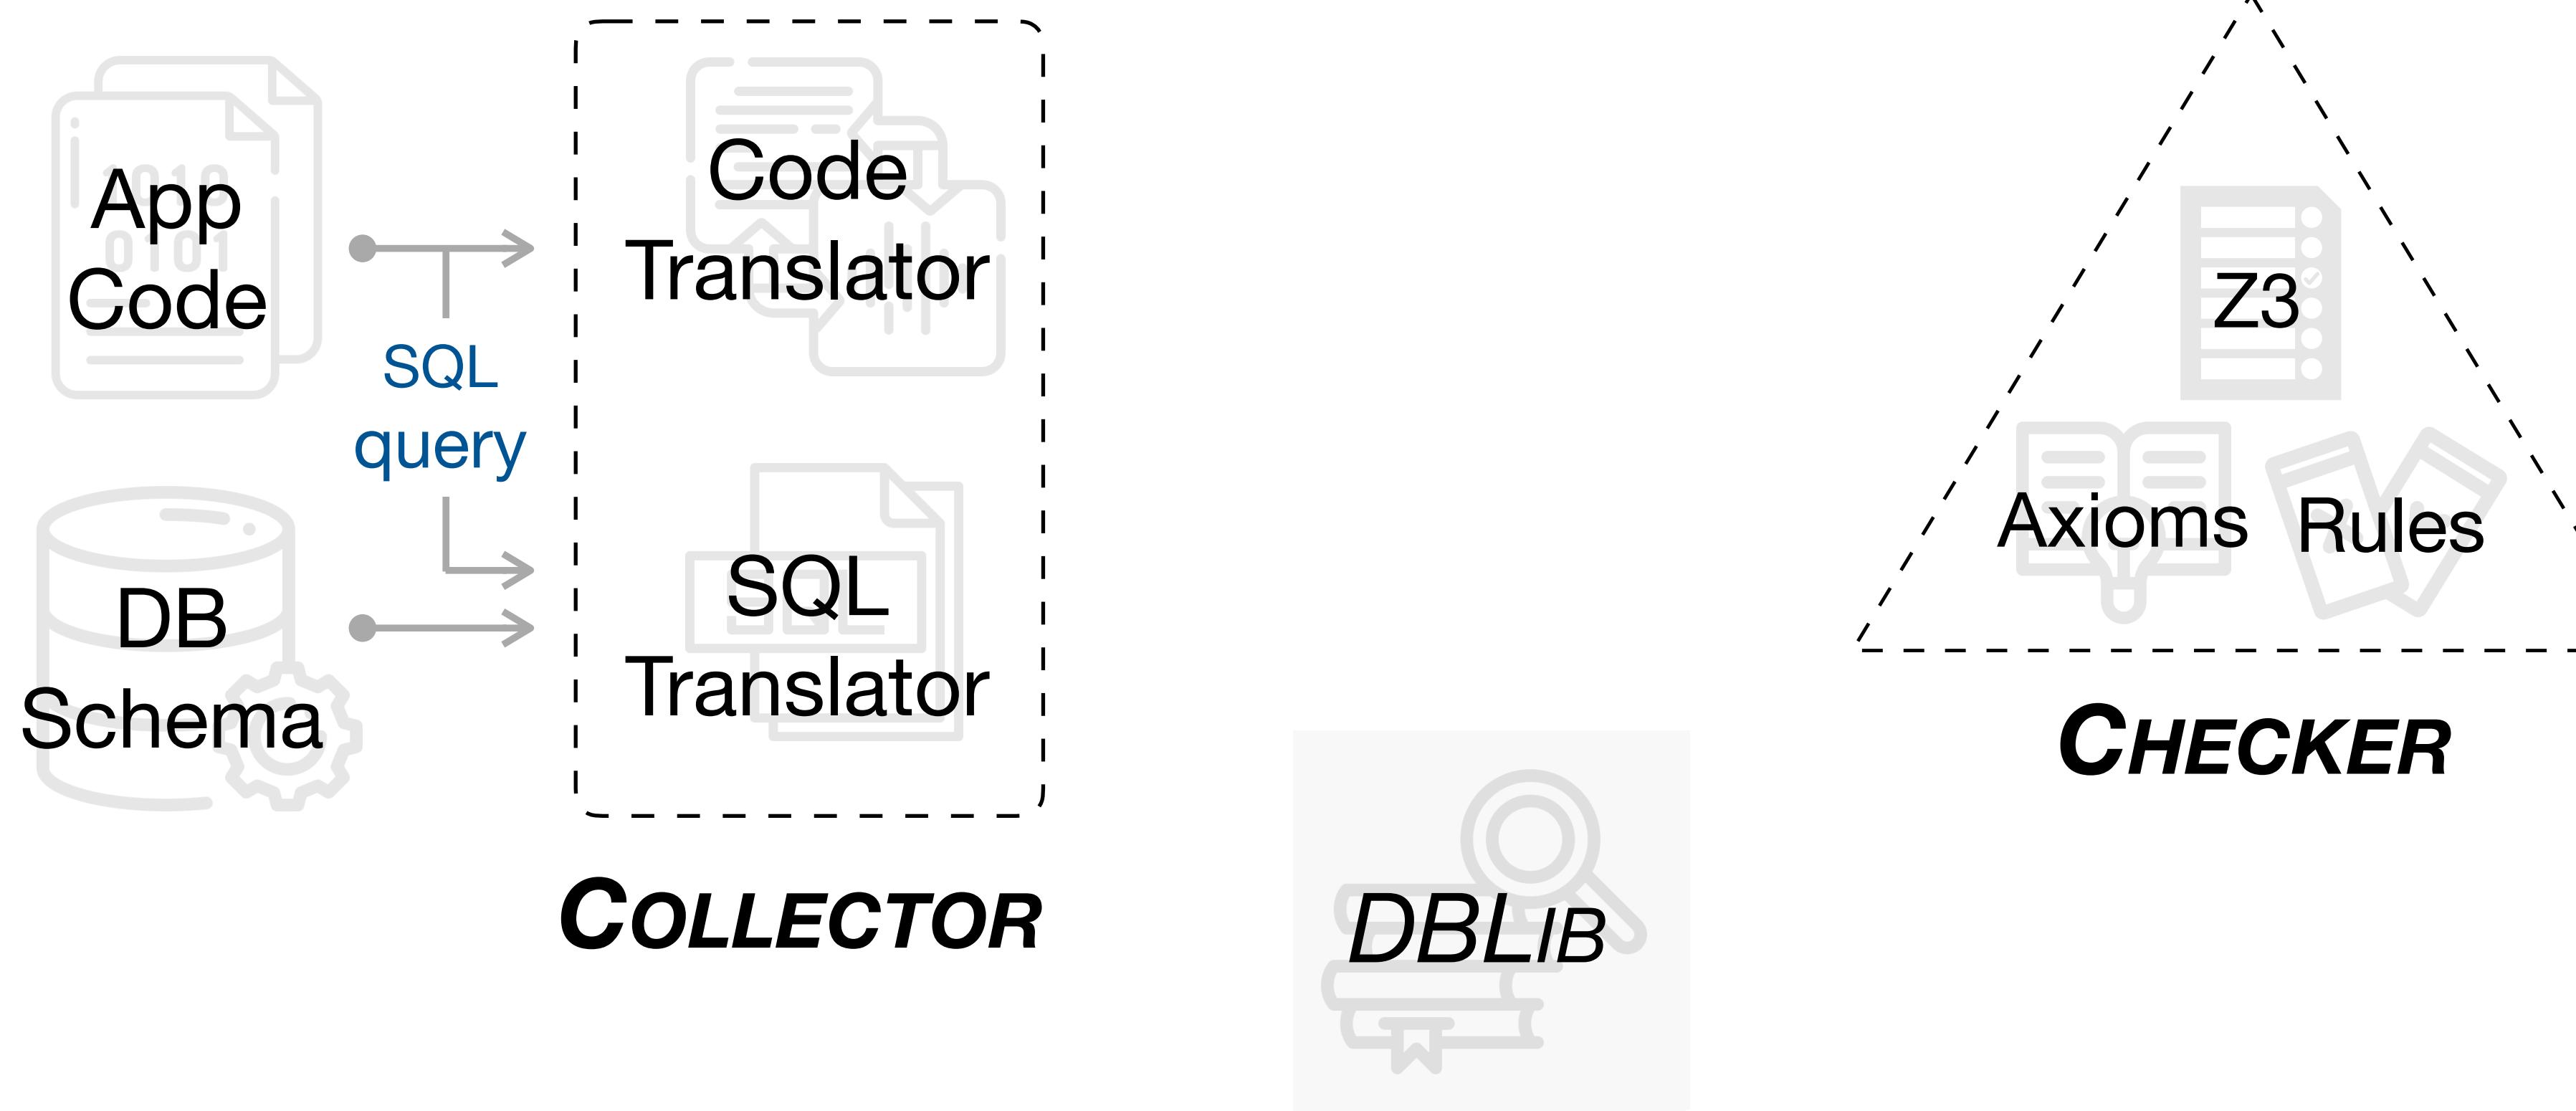


COLLECTOR

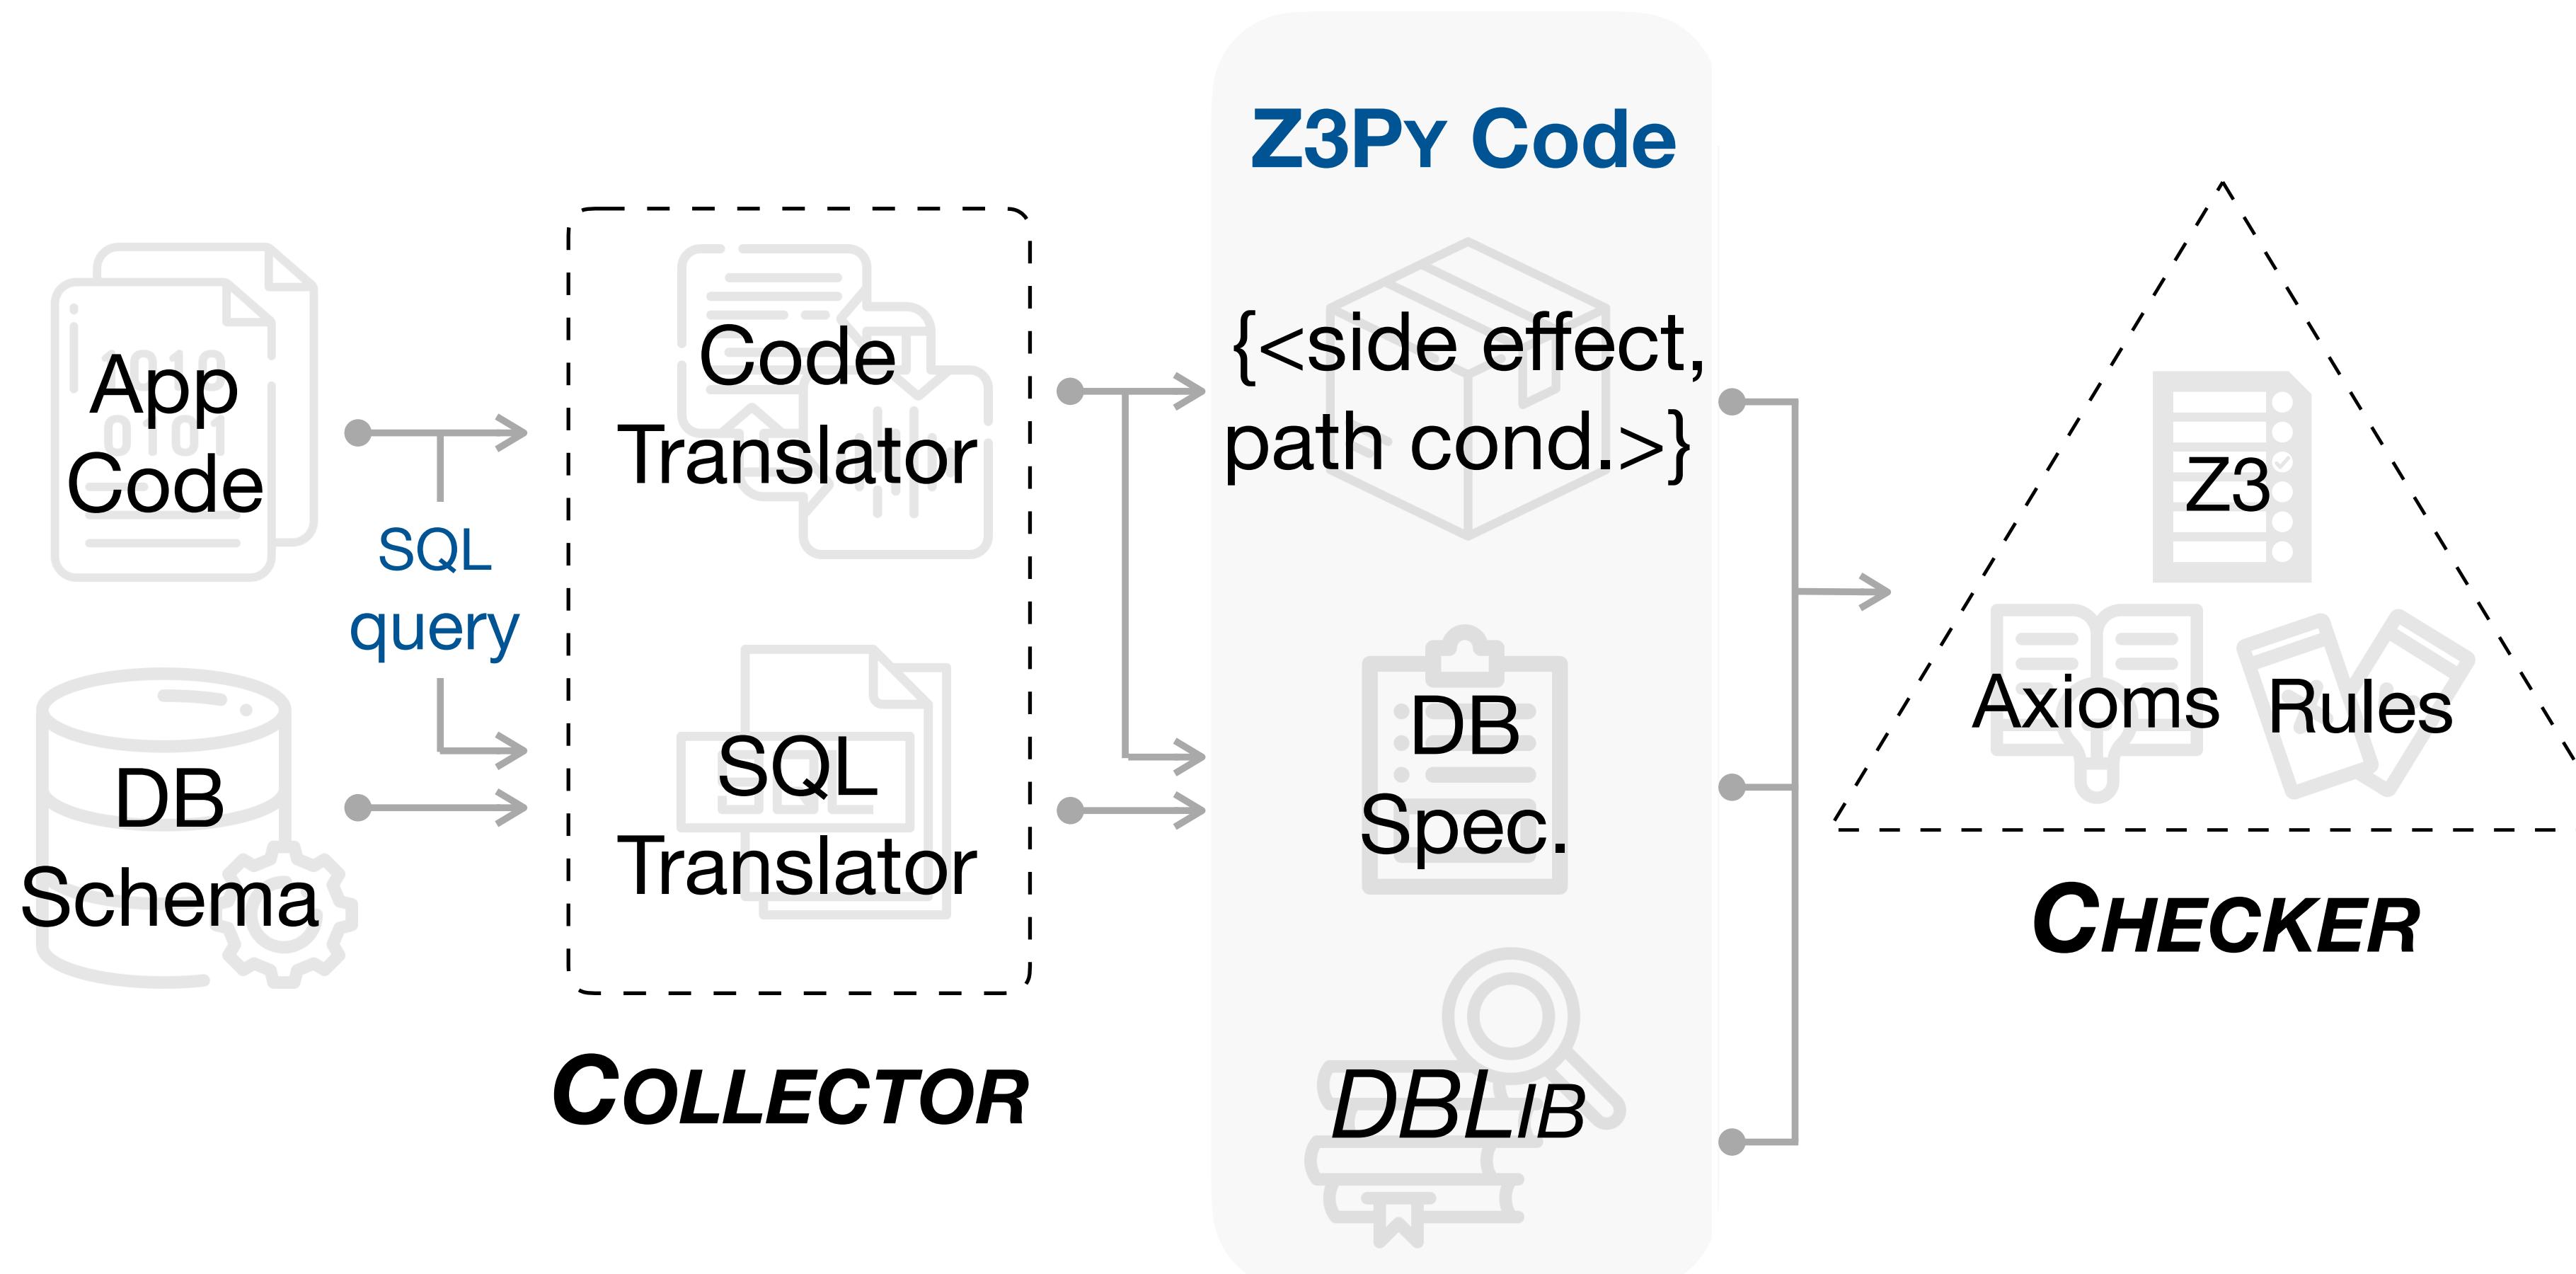


CHECKER

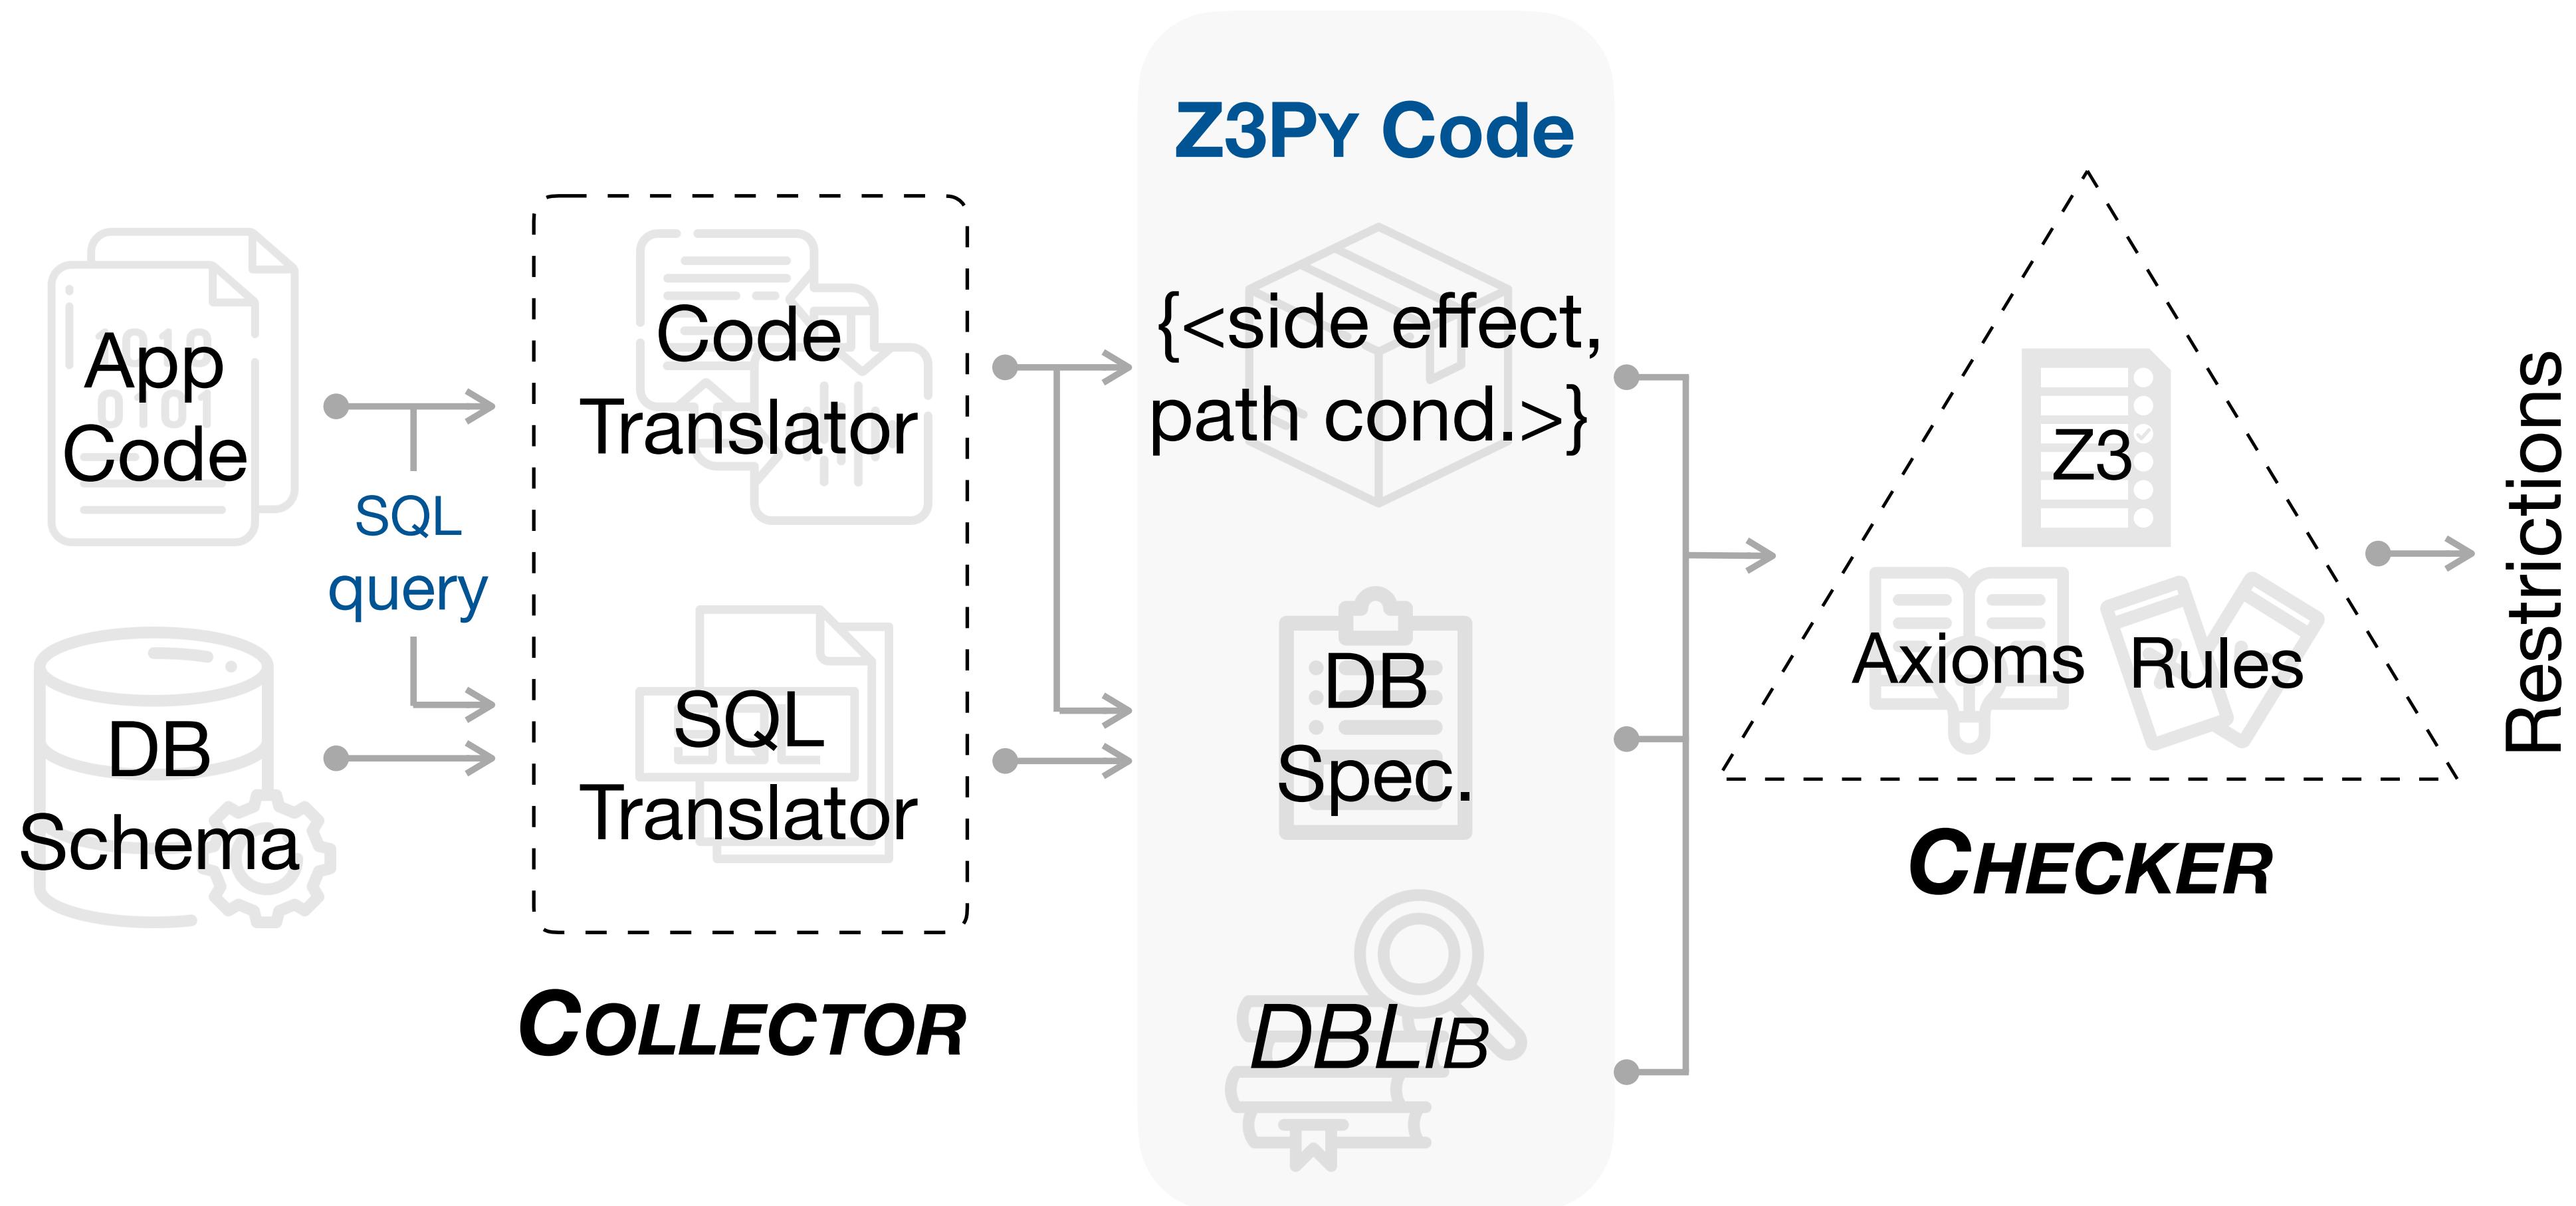
AutoGR --- Rigi



AutoGR --- Rigi



AutoGR --- Rigi

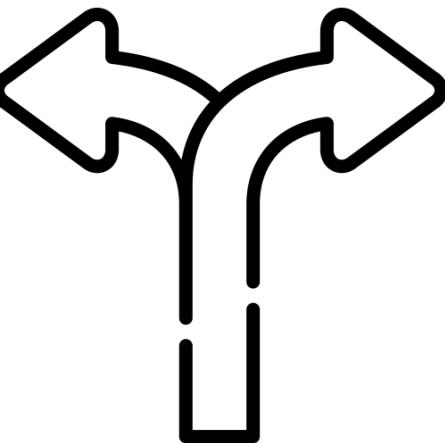


AutoGR --- Rigi --- DBLib

```
CREATE TABLE RSVN (
    R_C_ID  BIGINT NOT NULL ,
    R_F_ID  BIGINT NOT NULL ,
    R_SEAT   BIGINT NOT NULL , ...
PRIMARY KEY (R_C_ID , R_F_ID))
```

Table definition interface

- Primary key, foreign key
- Key with multiple fields



Z3Py

- IntSort, StringSort, BoolSort, RealSort, ...
- Extensional array
- ForAll, Implies, <=, ==, ...

```
K_RSVN = Datatype(...)
K_RSVN.declare(...,(R_C_ID , R_F_ID))
V_RSVN = Datatype(...)
V_RSVN.declare(...)
TABLE_RSVN = Array(..., K_RSVN , V_RSVN)
```

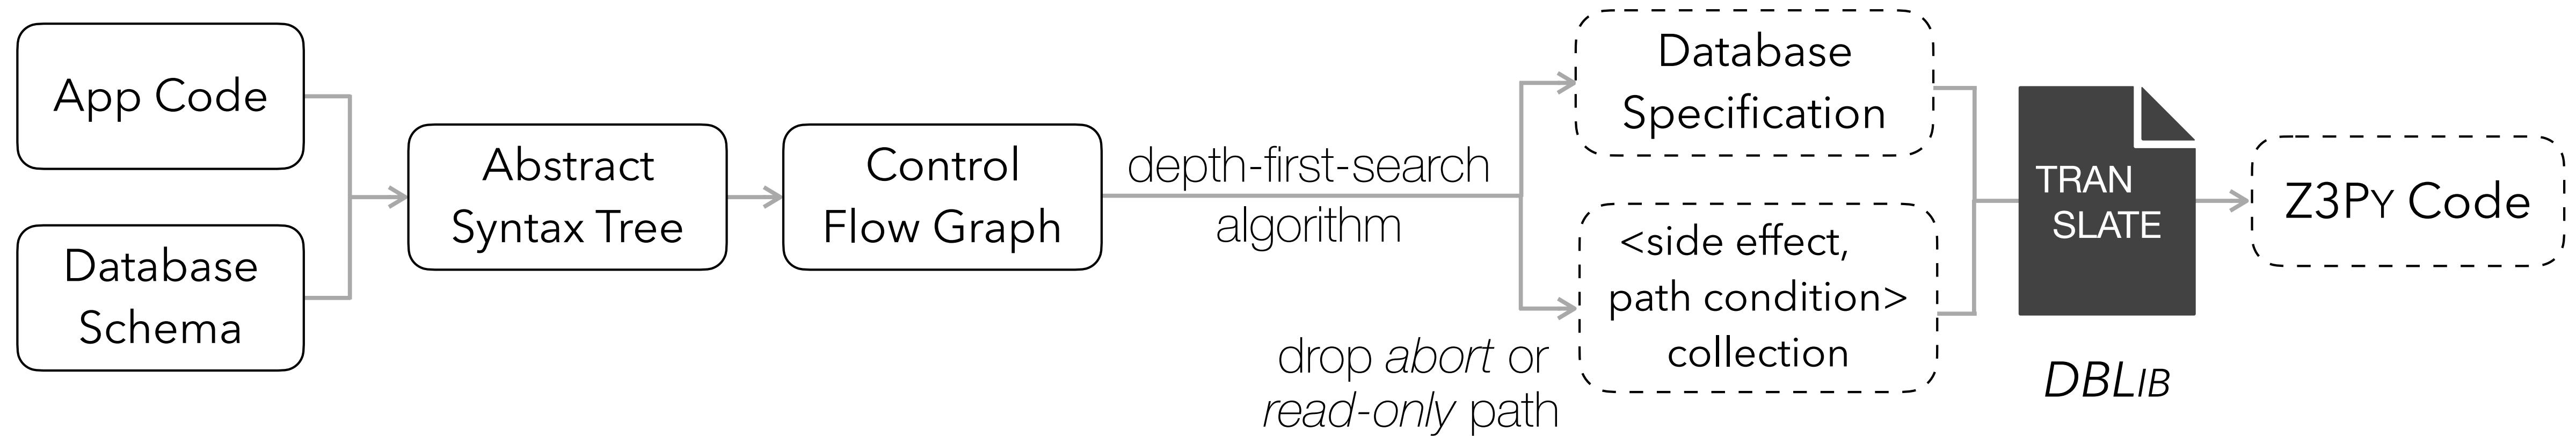
```
UPDATE CUSTOMER
SET C_BASE_AP_ID = aid ,
LOCATION = loc
WHERE C_ID = cid
```

SQL query interface

- Basic: select, update, ...
- Advanced: comparison, inner join, aggregation, ...

```
Store(TABLE_CUSTOMER , T_CUSTOMER.new(cid),
      V_CUSTOMER.new(V_CUSTOMER.BALANCE(
          Select(TABLE_CUSTOMER , T_CUSTOMER.new(cid)))
      ,aid ,... ,loc ,...))
```

AutoGR --- Rigi --- Collector



Optimizations:

- CRDTs support
 - e.g., "Last-Writer-Win (LWW)" strategy for merging concurrent updates.
- Uniqueness
 - Support database's AUTOINCREMENTAL feature.

AutoGR -- Rigi -- Checker

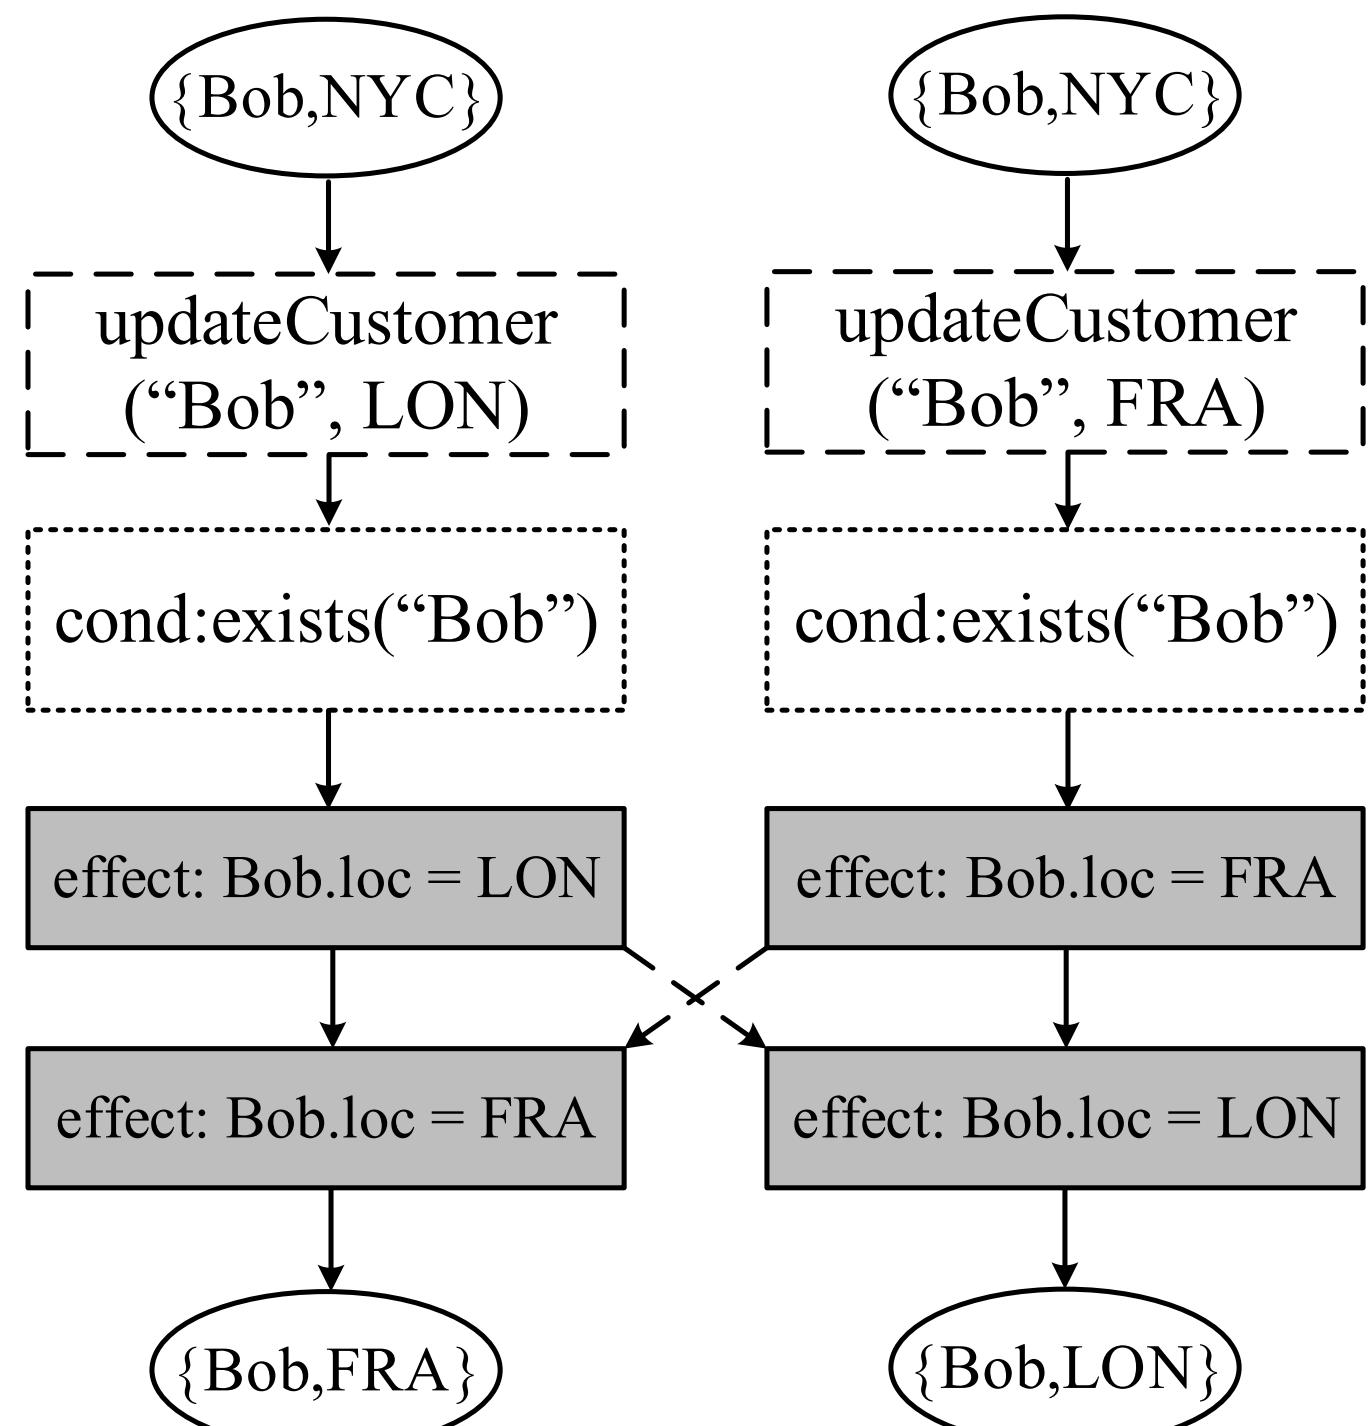
Commutativity check: Check the *commutativity of side-effects* to ensure the convergence of the system.

Semantics check: If the side-effect of operation A can be generated *without seeing the side-effect* of operation B, then it must be able to be generated *when seeing the side-effect* of operation B.

AutoGR -- Rigi -- Checker

Commutativity check: Check the *commutativity of side-effects* to ensure the convergence of the system.

Semantics check: If the side-effect of operation A can be generated *without seeing the side-effect* of operation B, then it must be able to be generated *when seeing the side-effect* of operation B.

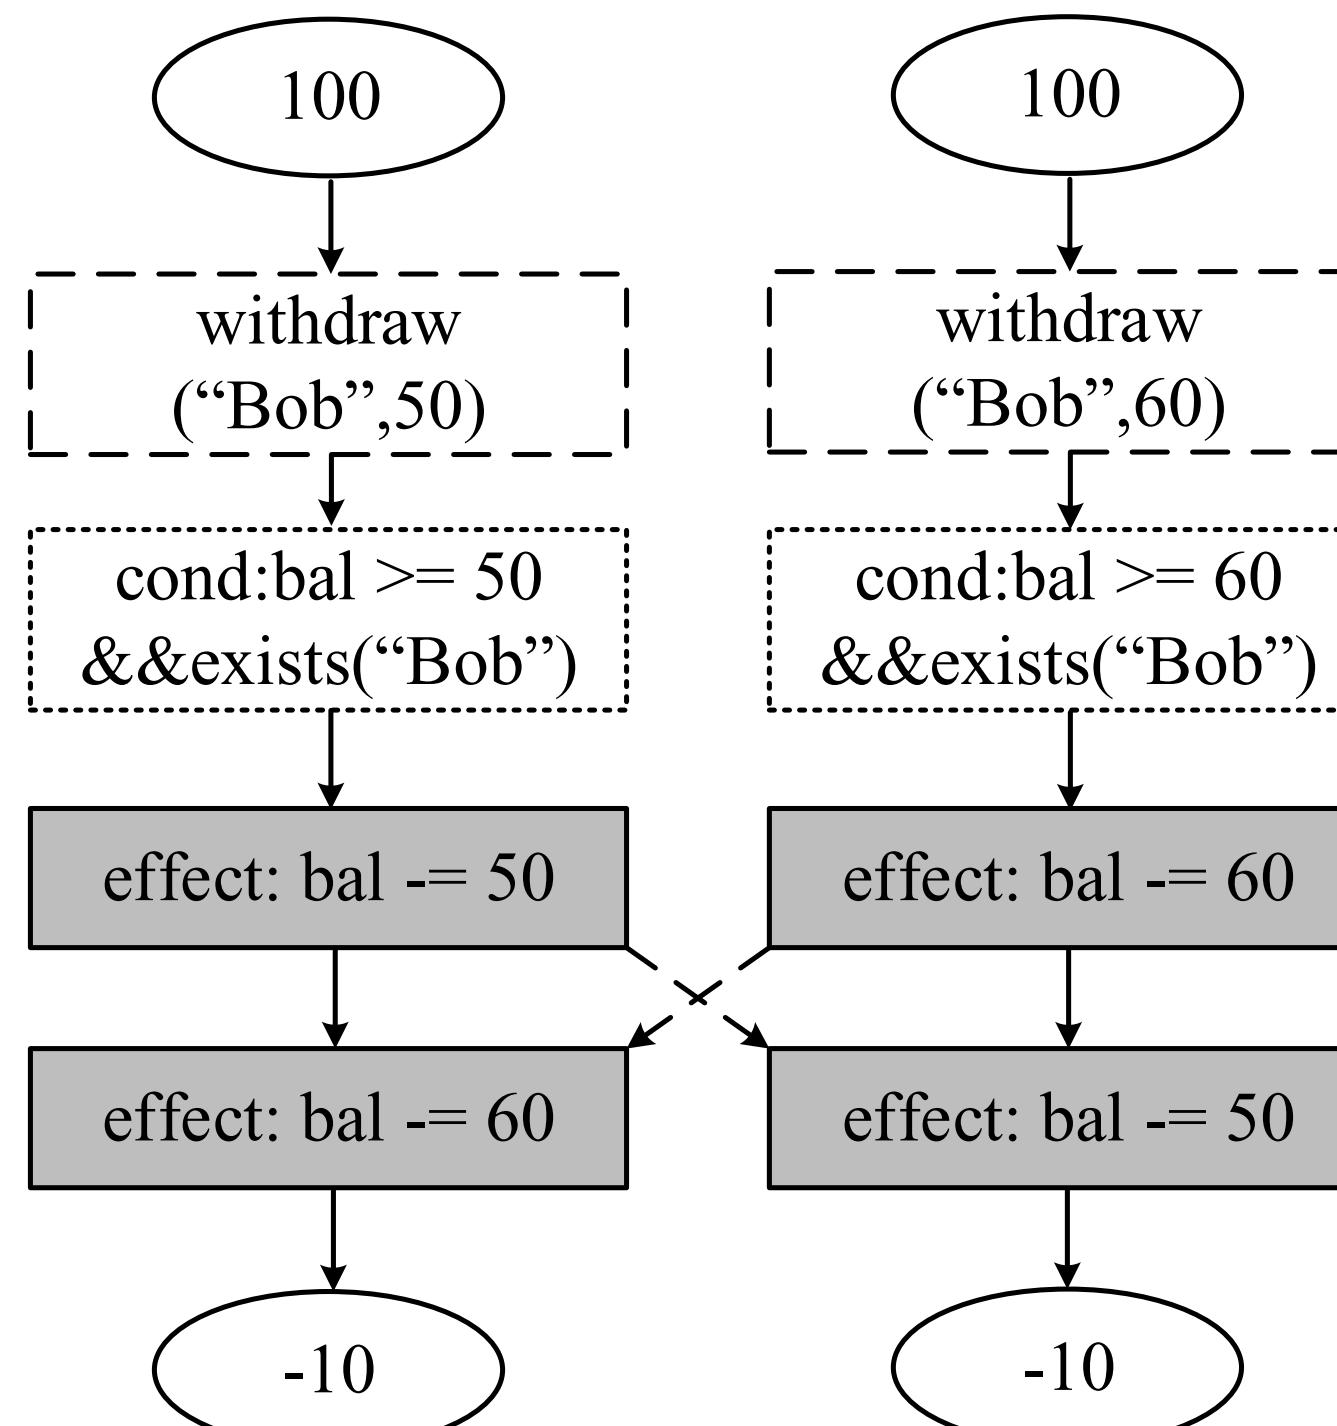


(a) Violating execution 1

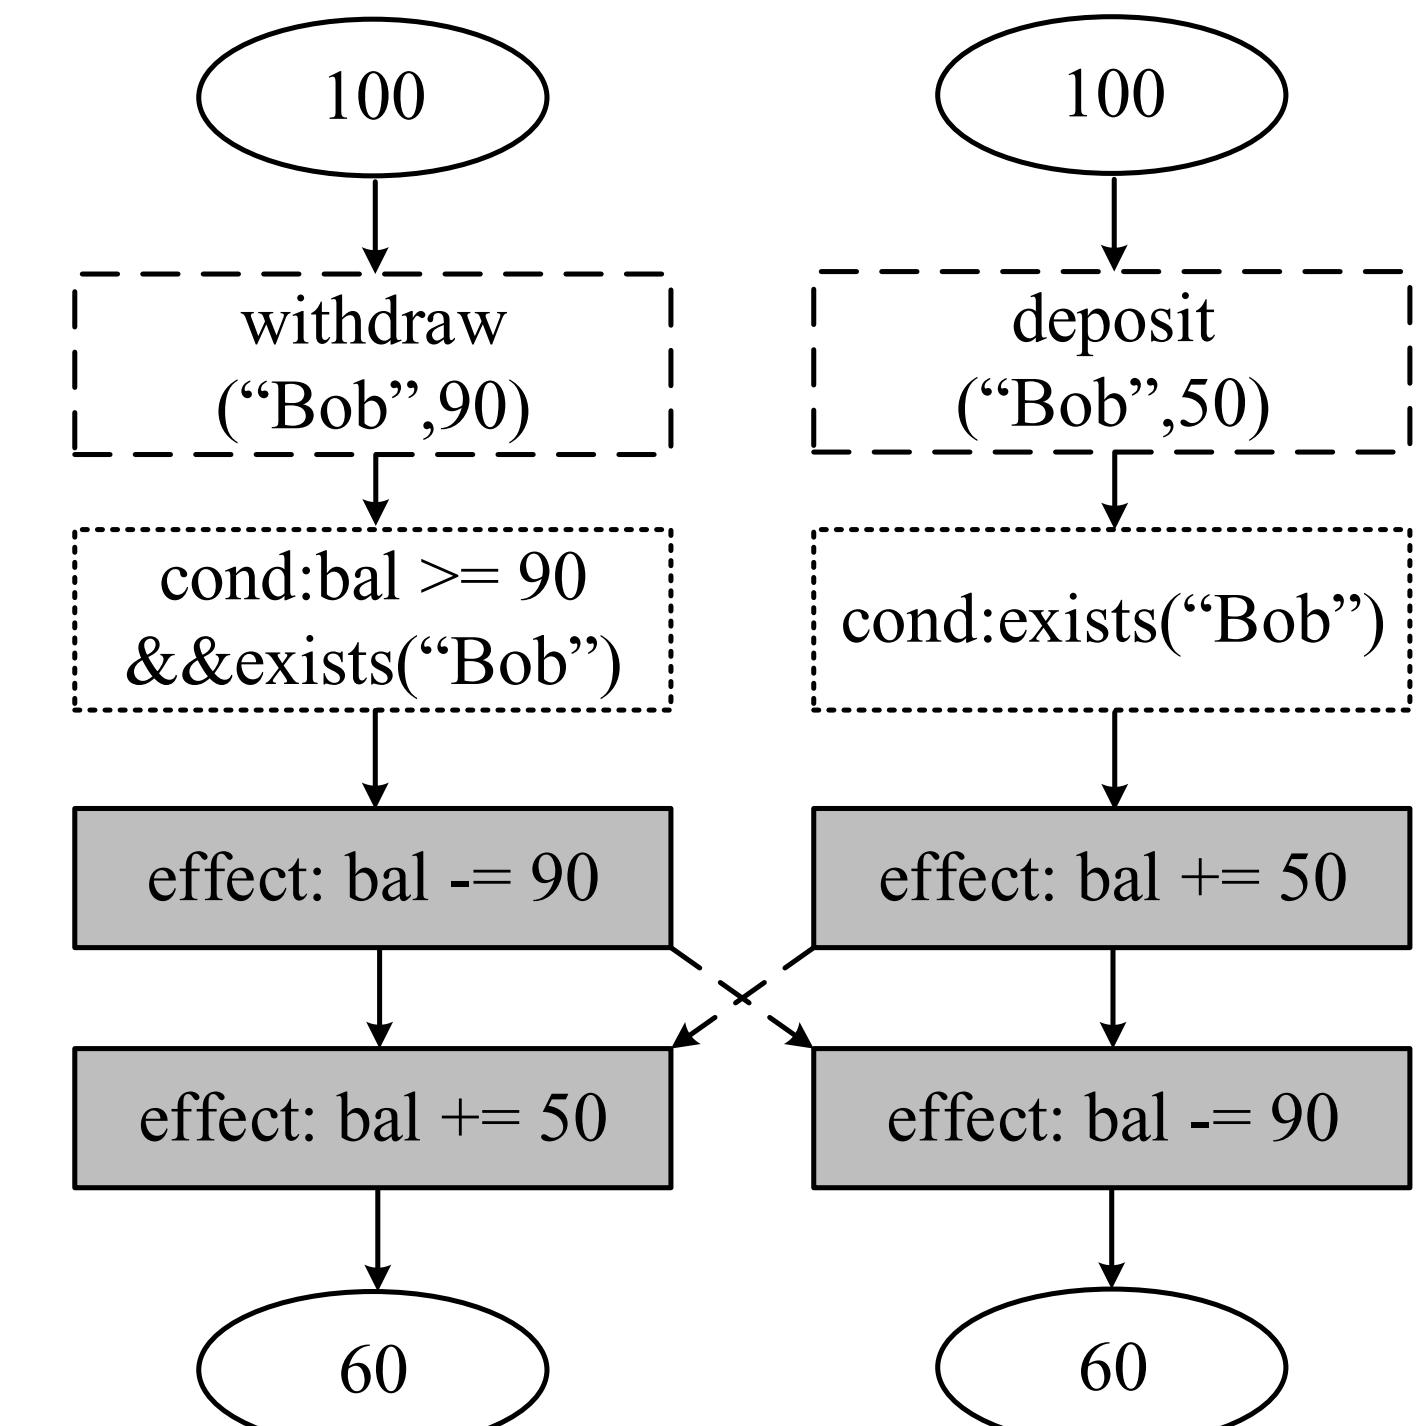
AutoGR --- Rigi --- Checker

Commutativity check: Check the *commutativity of side-effects* to ensure the convergence of the system.

Semantics check: If the side-effect of operation A can be generated *without seeing the side-effect* of operation B, then it must be able to be generated *when seeing the side-effect* of operation B.



(b) Violating execution 2

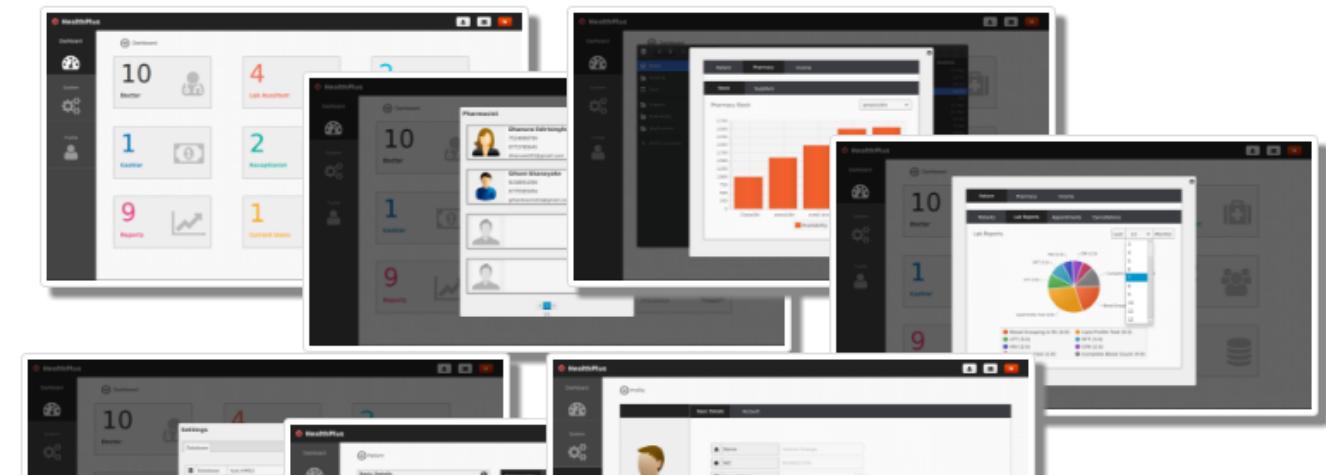


(c) Correct execution.

Case Study

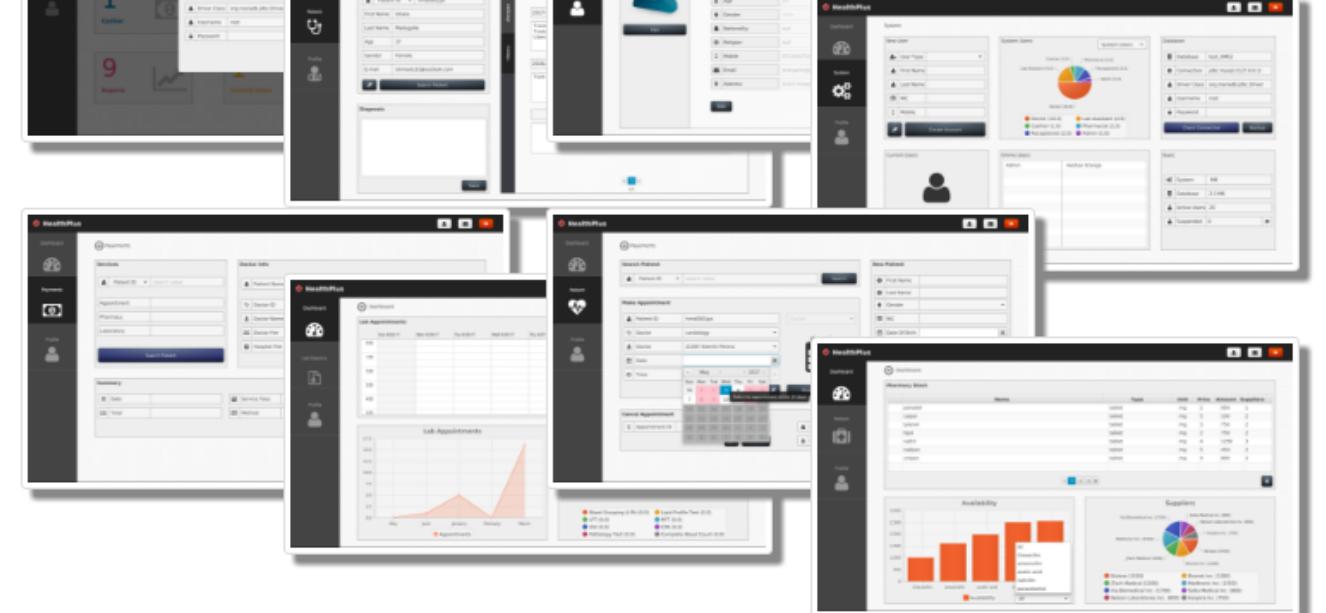
- *SmallBank* (codebase 2.5k, 5 transactions):

Simulating an online banking system.



- *RUBiS* (codebase 9.8k, 16 transactions):

An eBay-like online auction website.

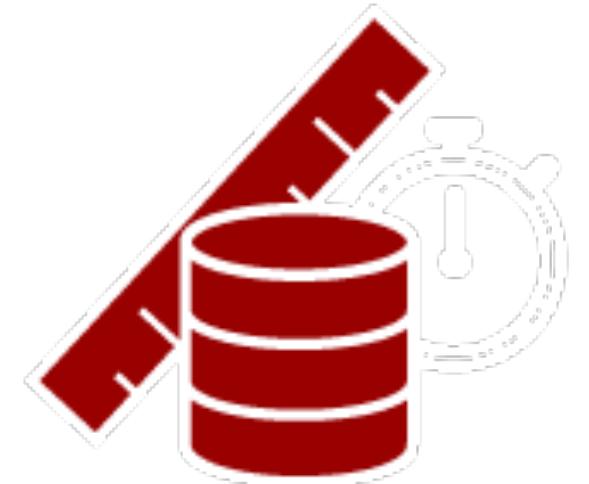


- *SeatsReservation* (codebase 5.0k, 6 transactions):

An electronic airline ticketing service.

- *HealthPlus* (codebase 15.7k, 157 transactions):

A real-world deployable management system for health care facility.



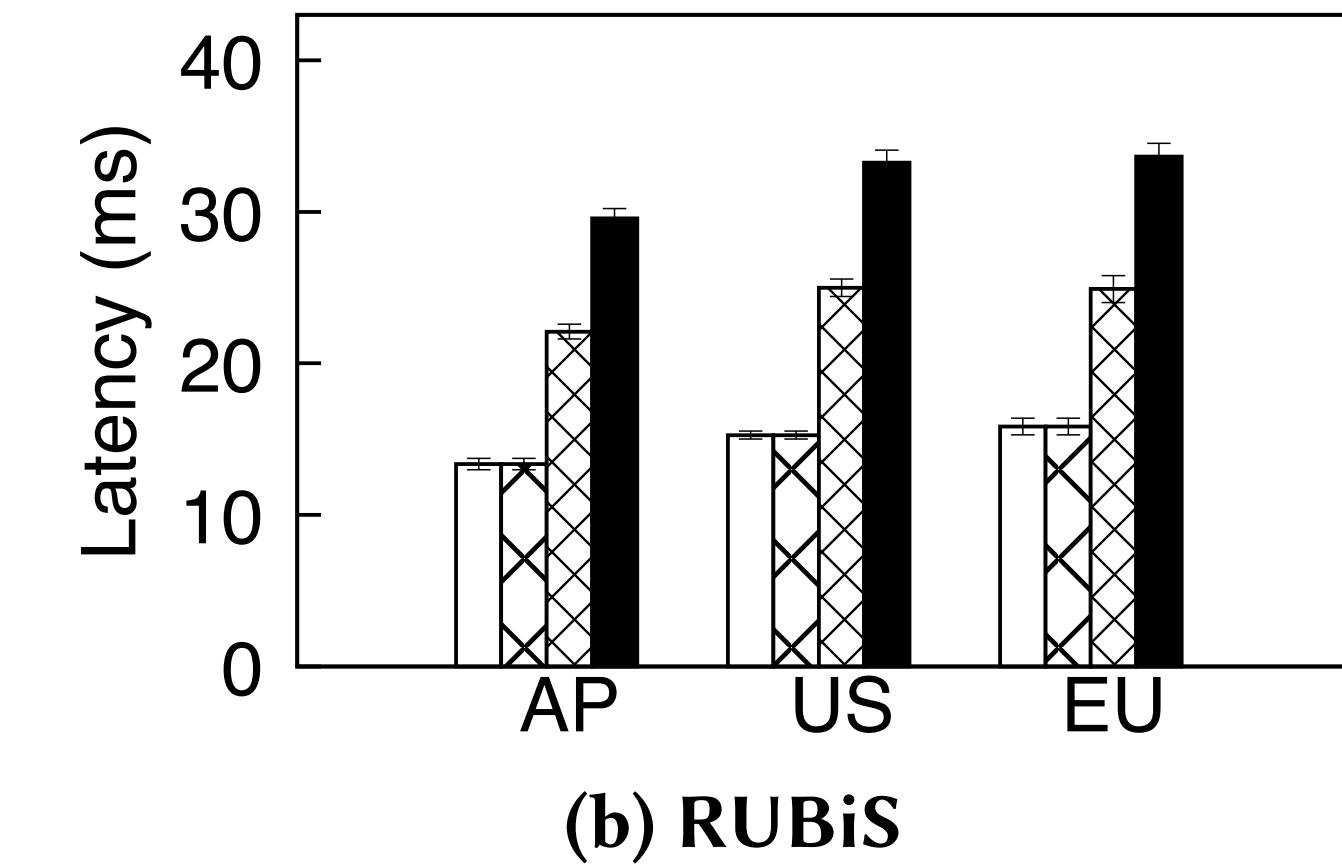
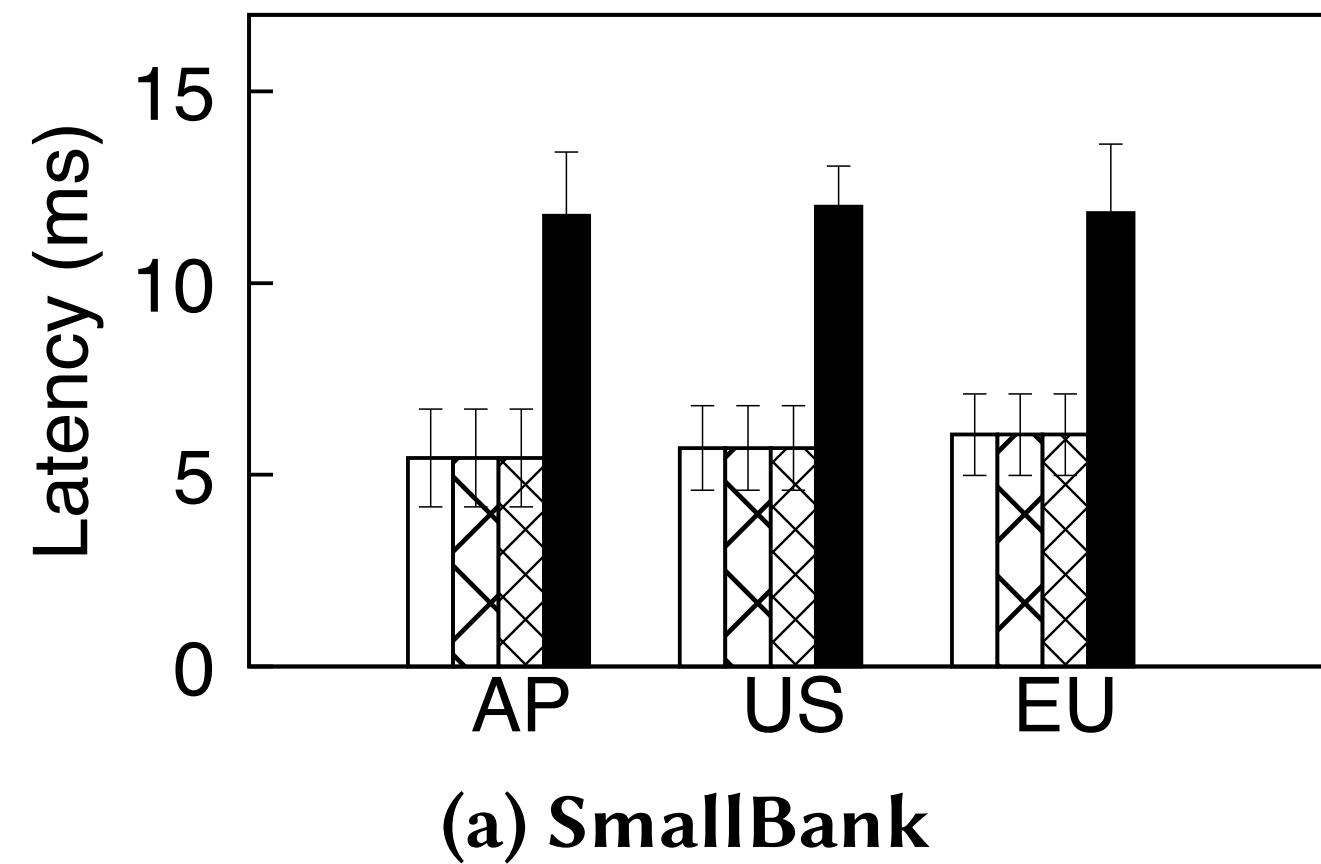
Case Study — Static analysis

	Lines of Z3Py Code Generated by Rigi			Analysis Cost	Restriction Rate (Normal / Opt)
	<u>Database</u>	<u>Path condition</u>	<u>Side effect</u>		
SmallBank	29	38	119	~ 24s	20% / 20%
RUBiS	113	62	191	~ 3.4min	23% / 9%
Seats	267	65	207	~ 6.7min	39% / 31%
HealthPlus	524	1113	1387	~ 1.5h (~ 7.7min for 16 threads)	2.9% / 1.4%

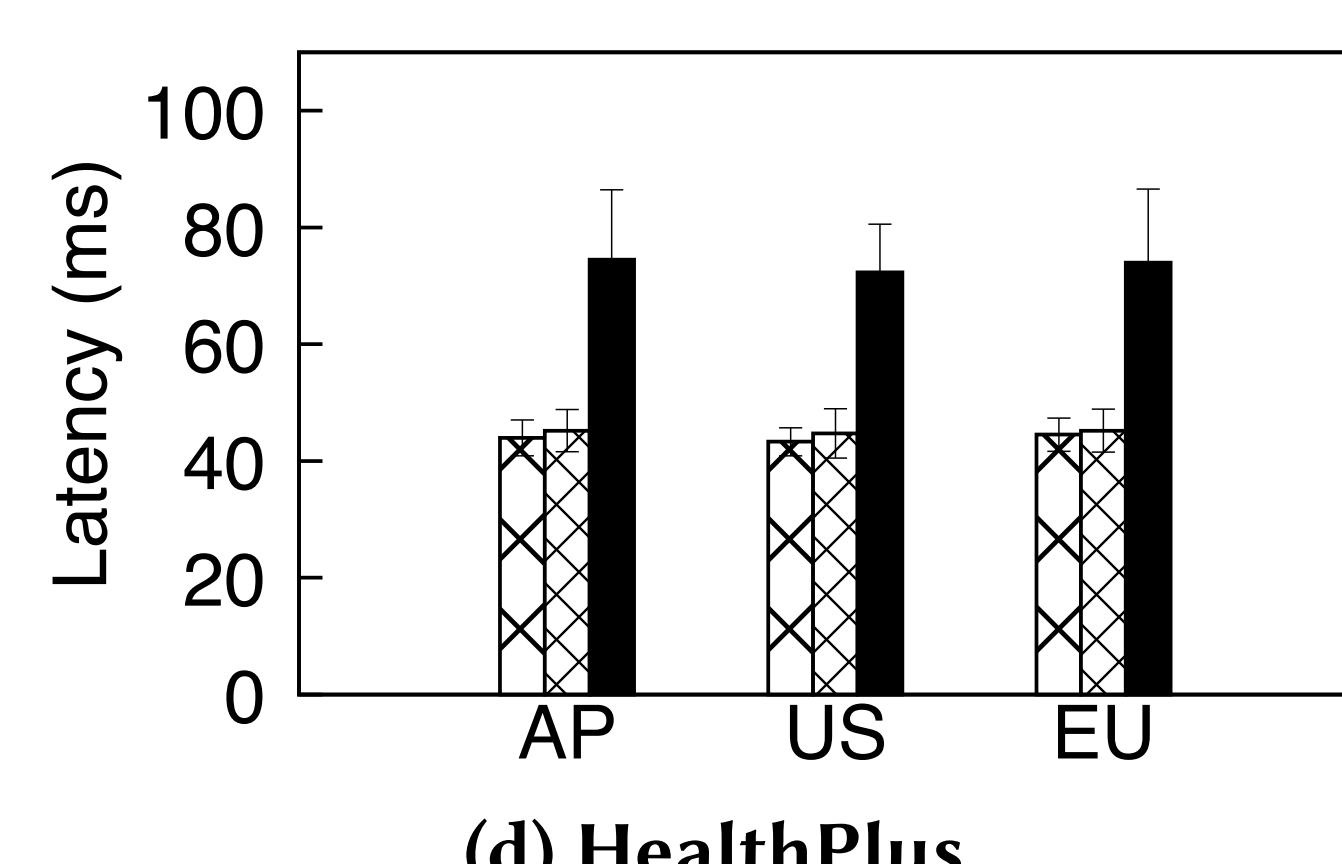
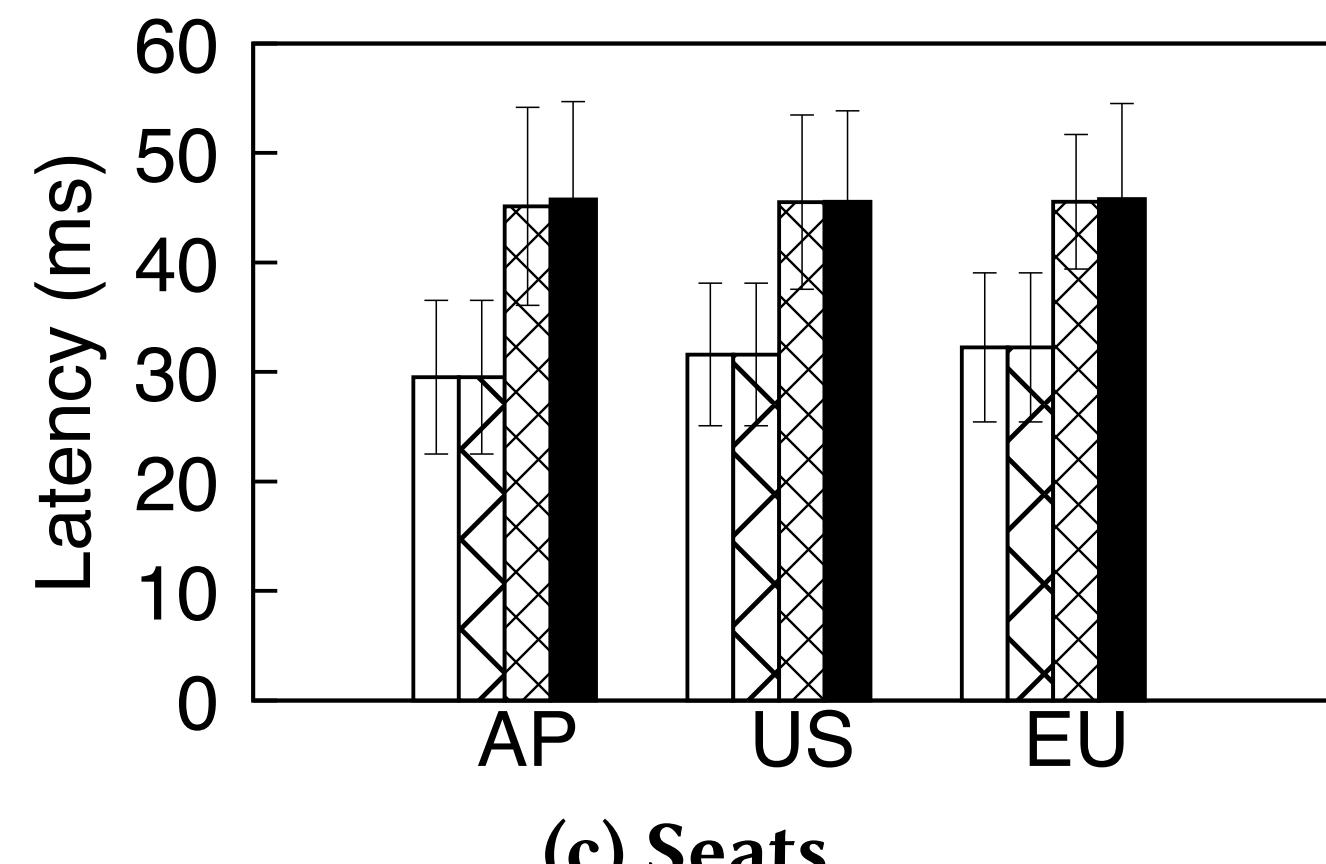
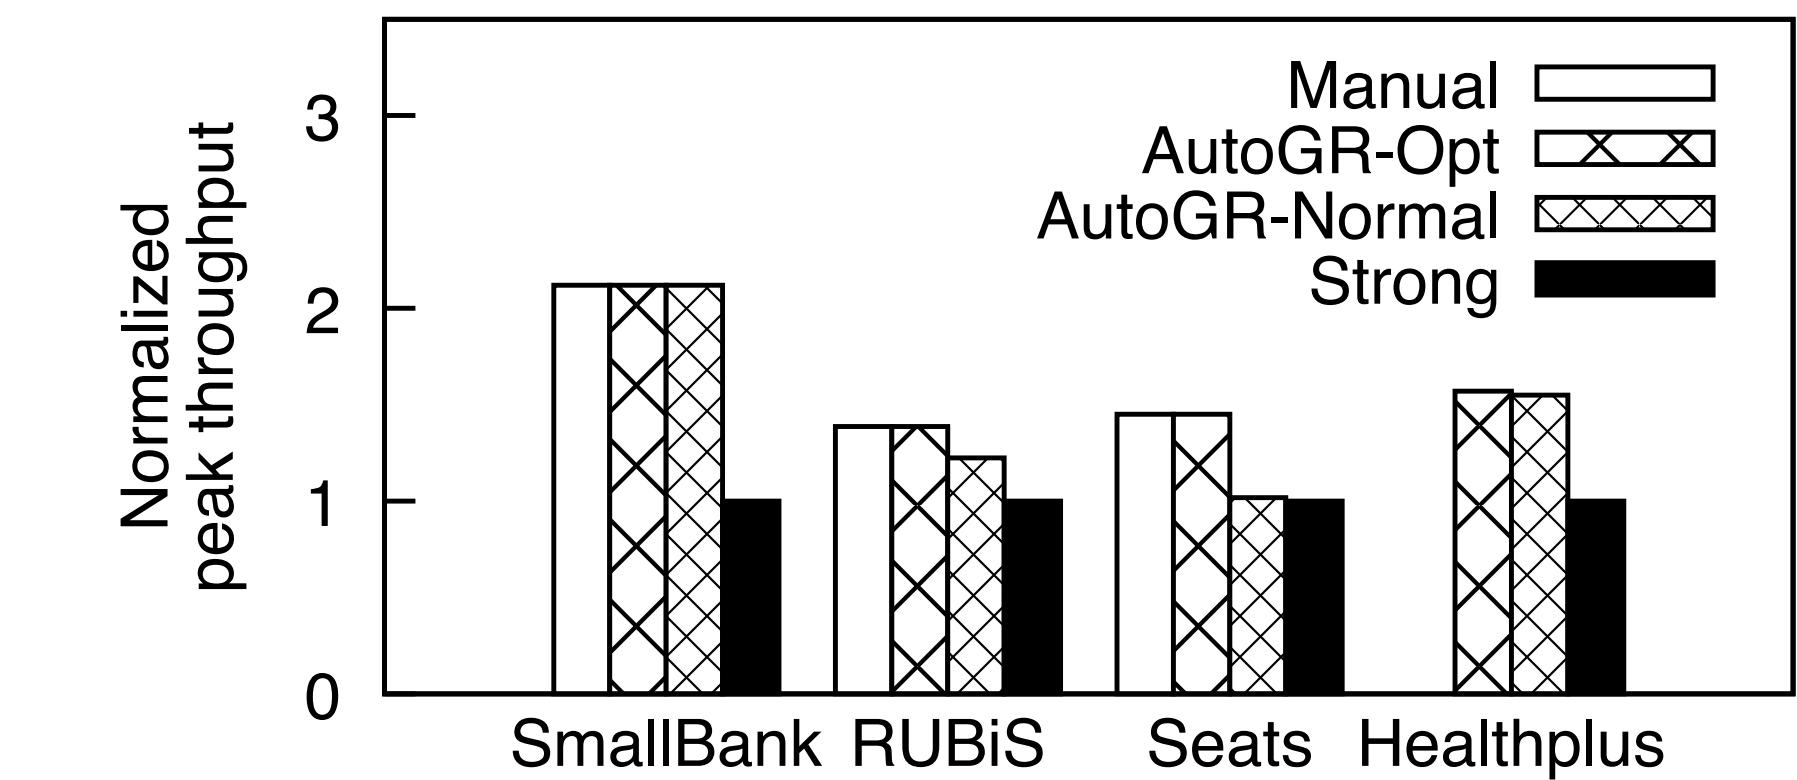
Considering that the analysis is a one-time and offline job, the cost is moderate.

Case Study — Geo-replication

Average latency perceived by users



Normalized peak throughput numbers



Manual

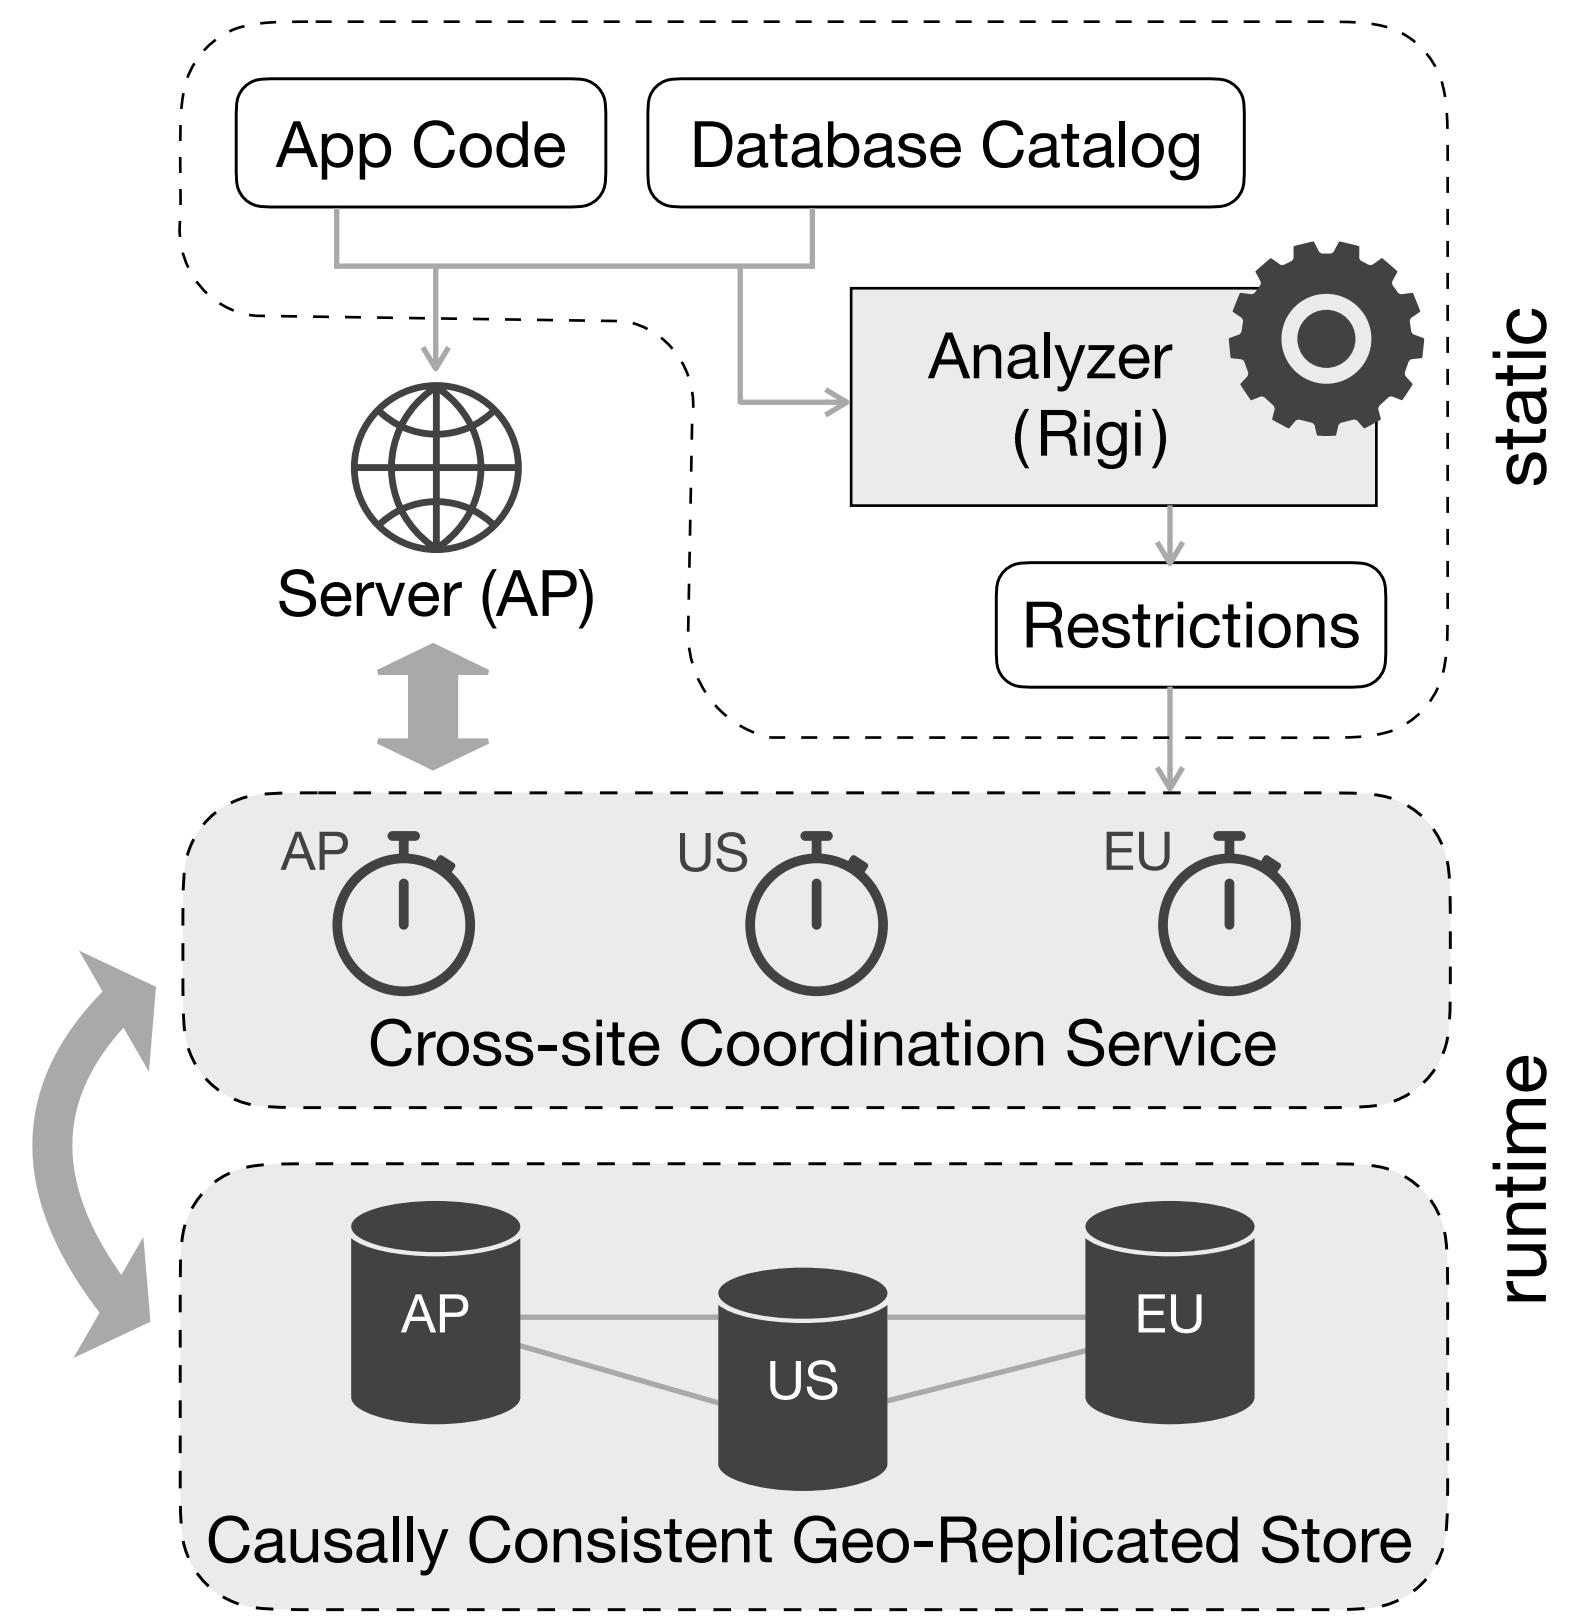
AutoGR-Opt

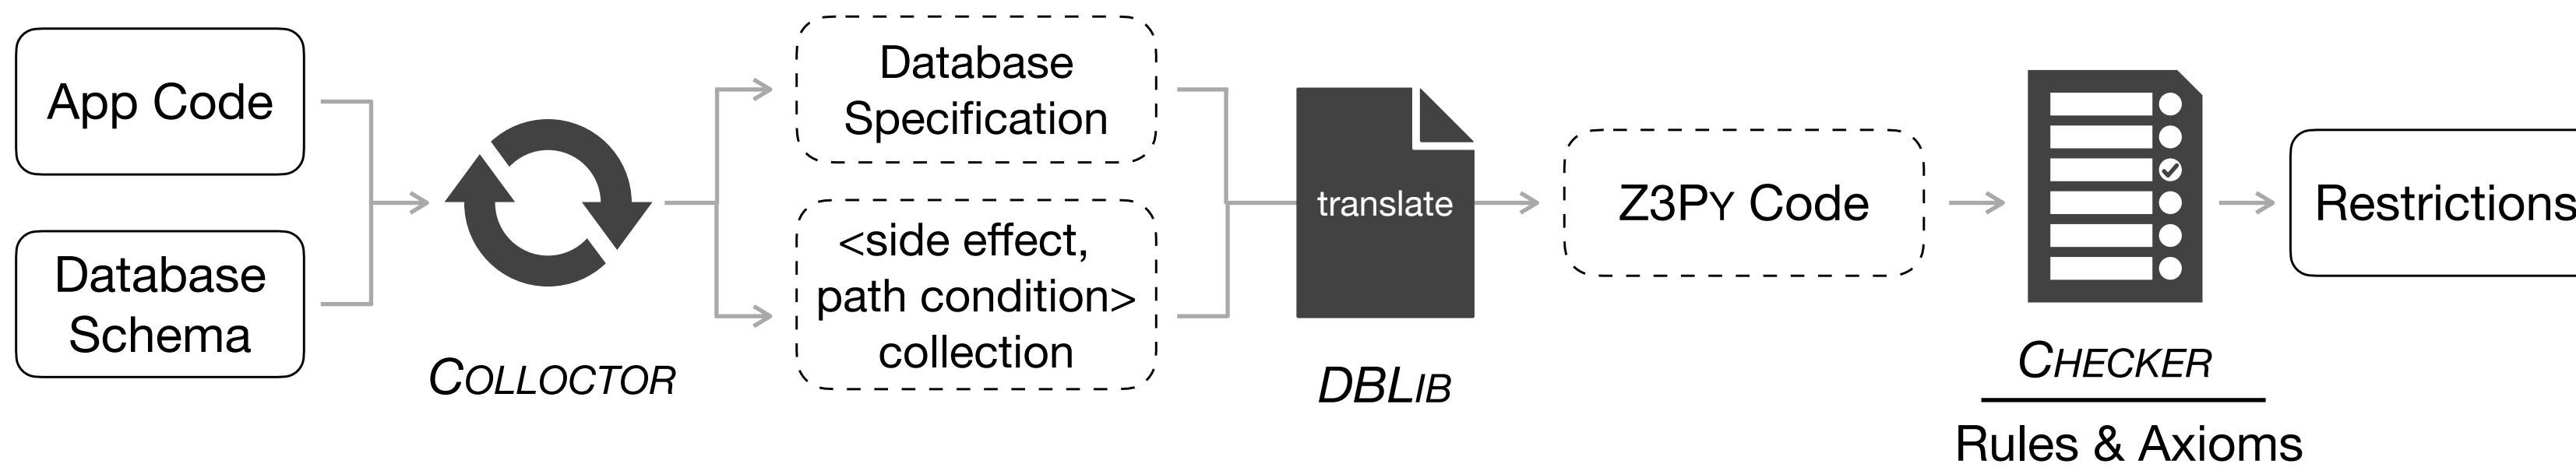
AutoGR-Normal

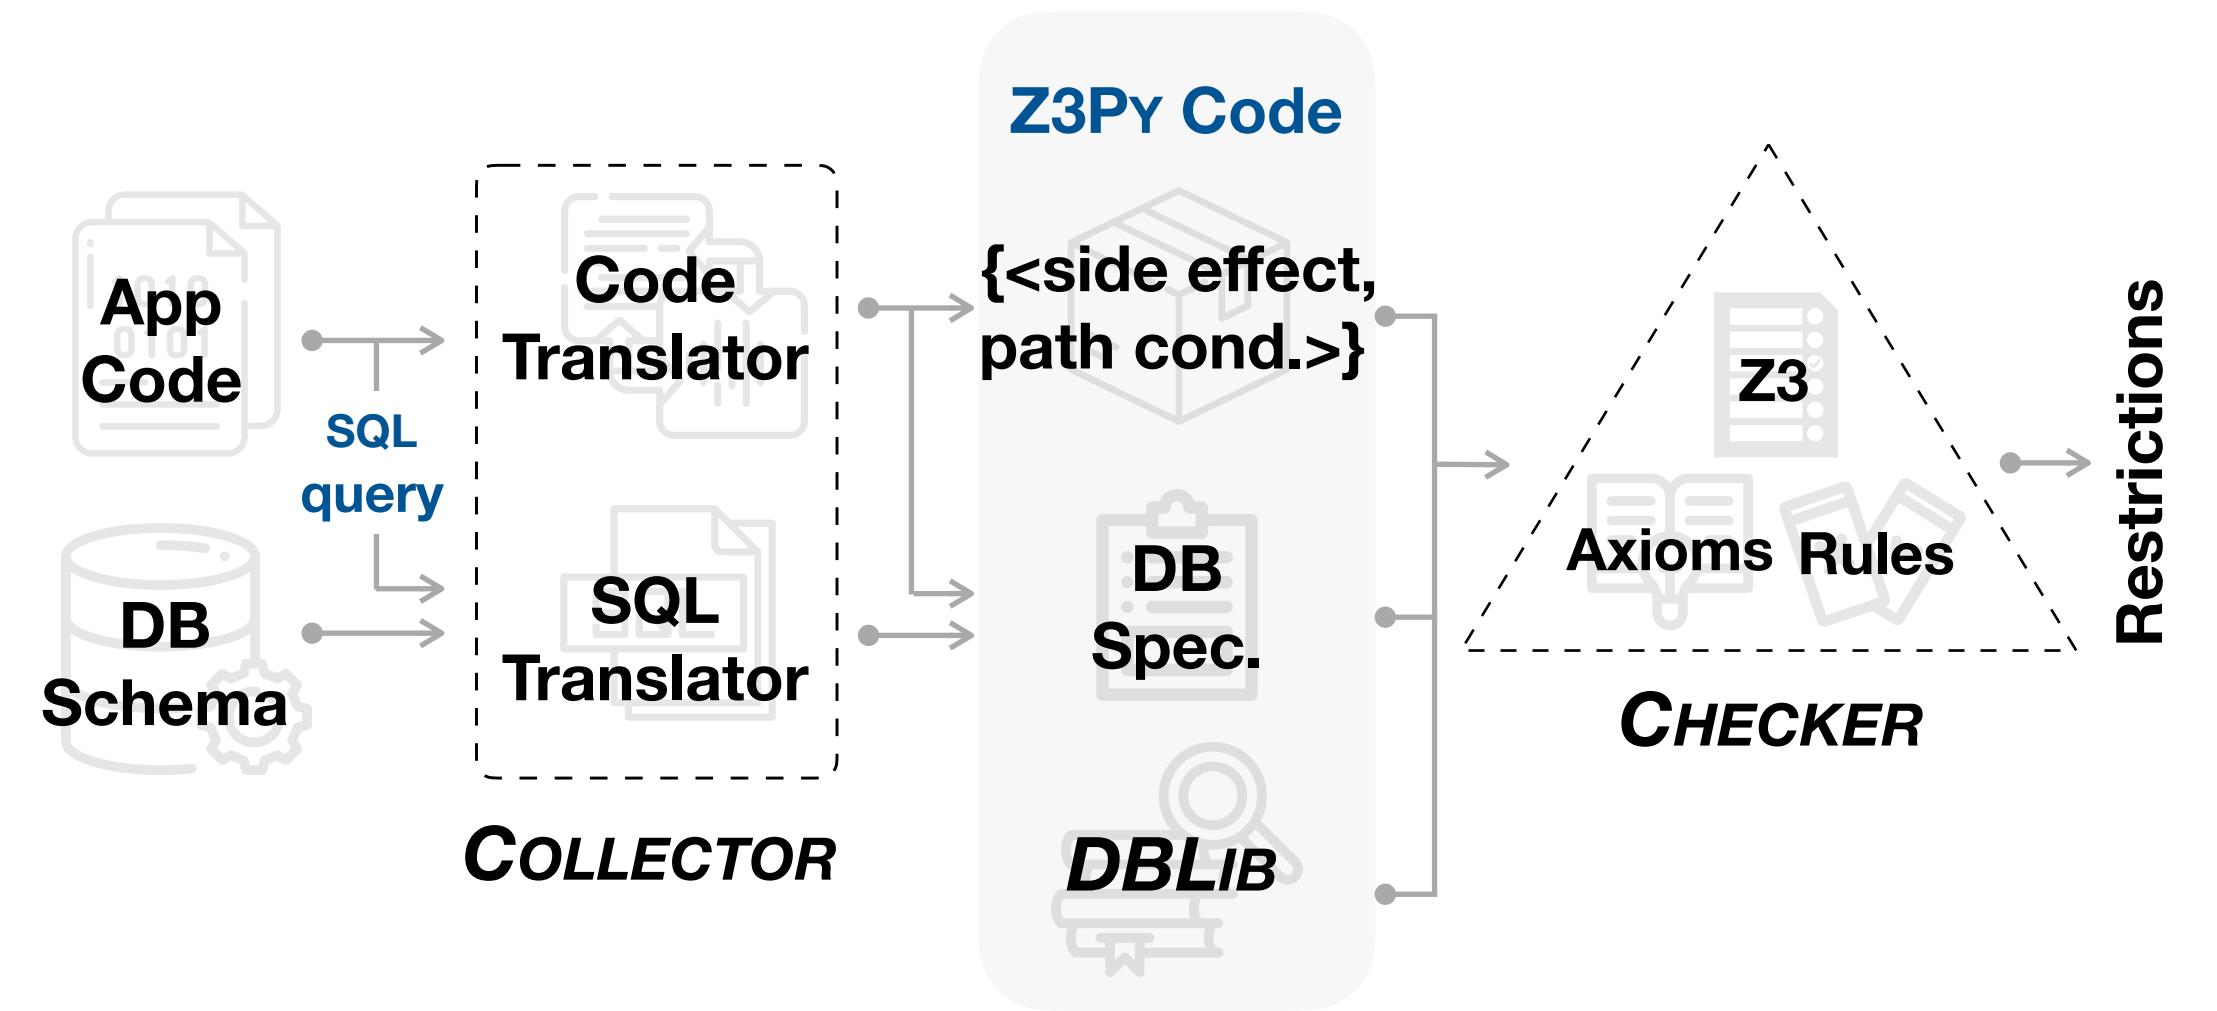
Strong

- Compare with human-intervention-free automated approaches: reduces up to 61.8% latency and achieves up to 2.12× higher peak throughput.
- Compare with manual analysis approaches: quickly enable the geo-replication feature with zero human intervention while offering similarly low latency and high throughput.

Thank You !







Manual 

AutoGR-Opt 

AutoGR-Normal 

Strong 

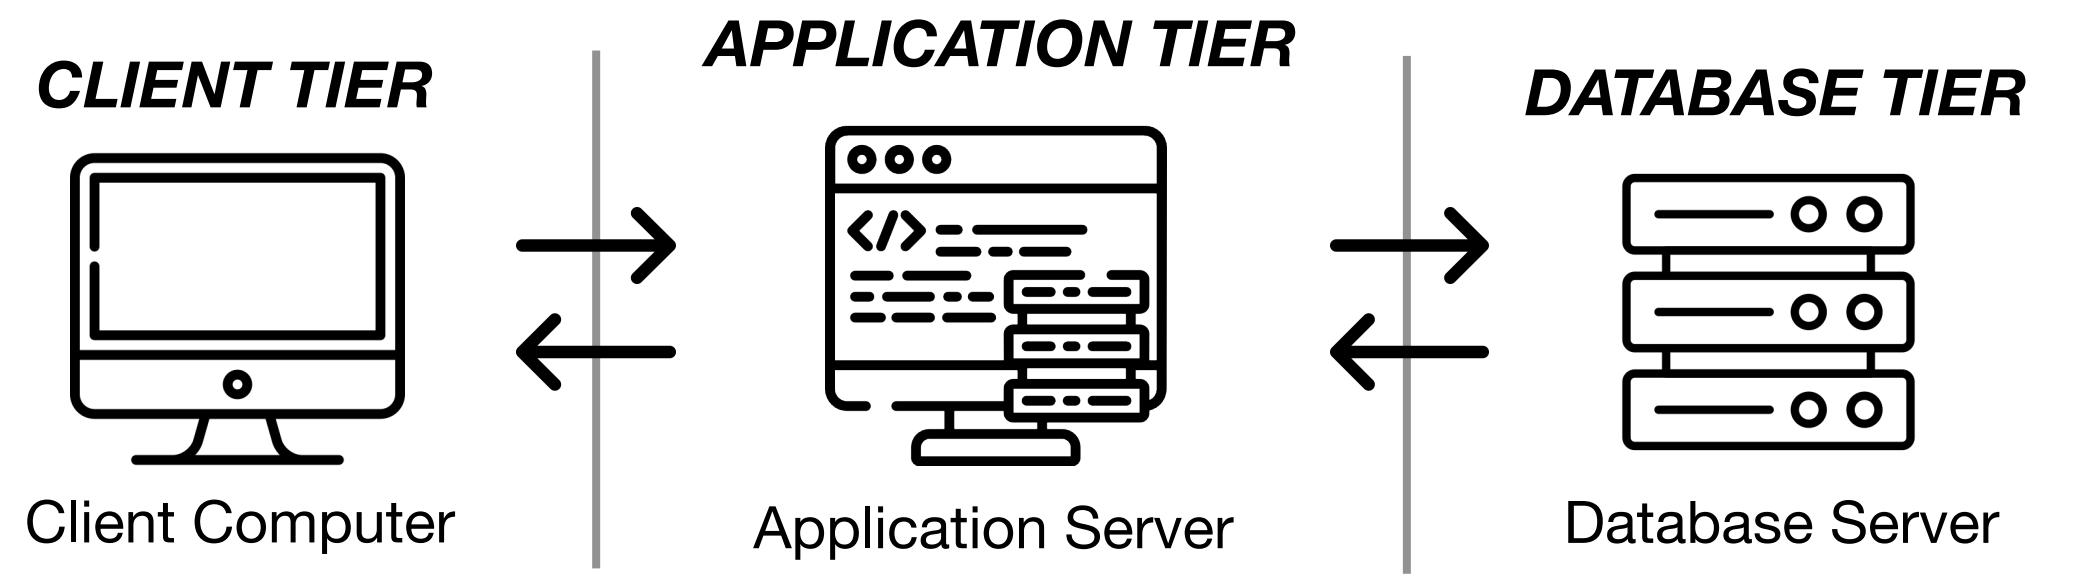
Single-Site-Strong □

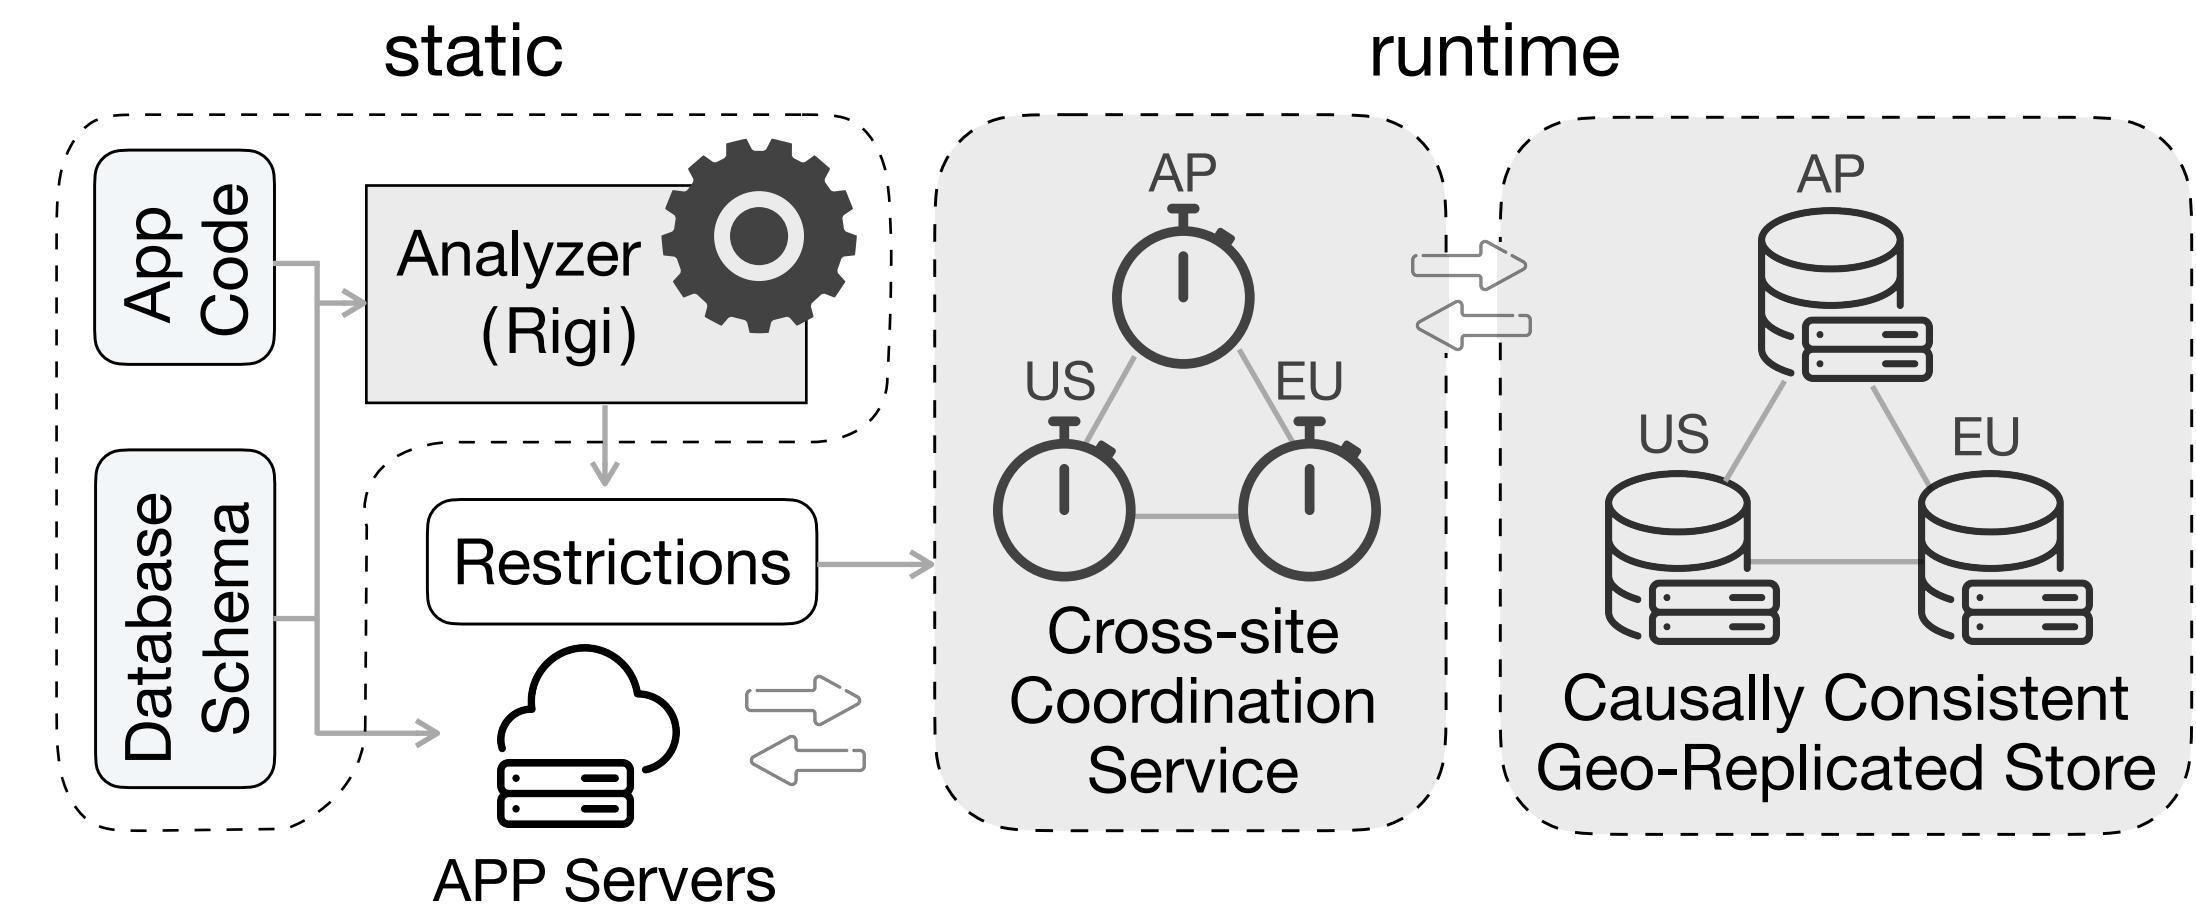
3-Site-AutoGR-Optimal ✕

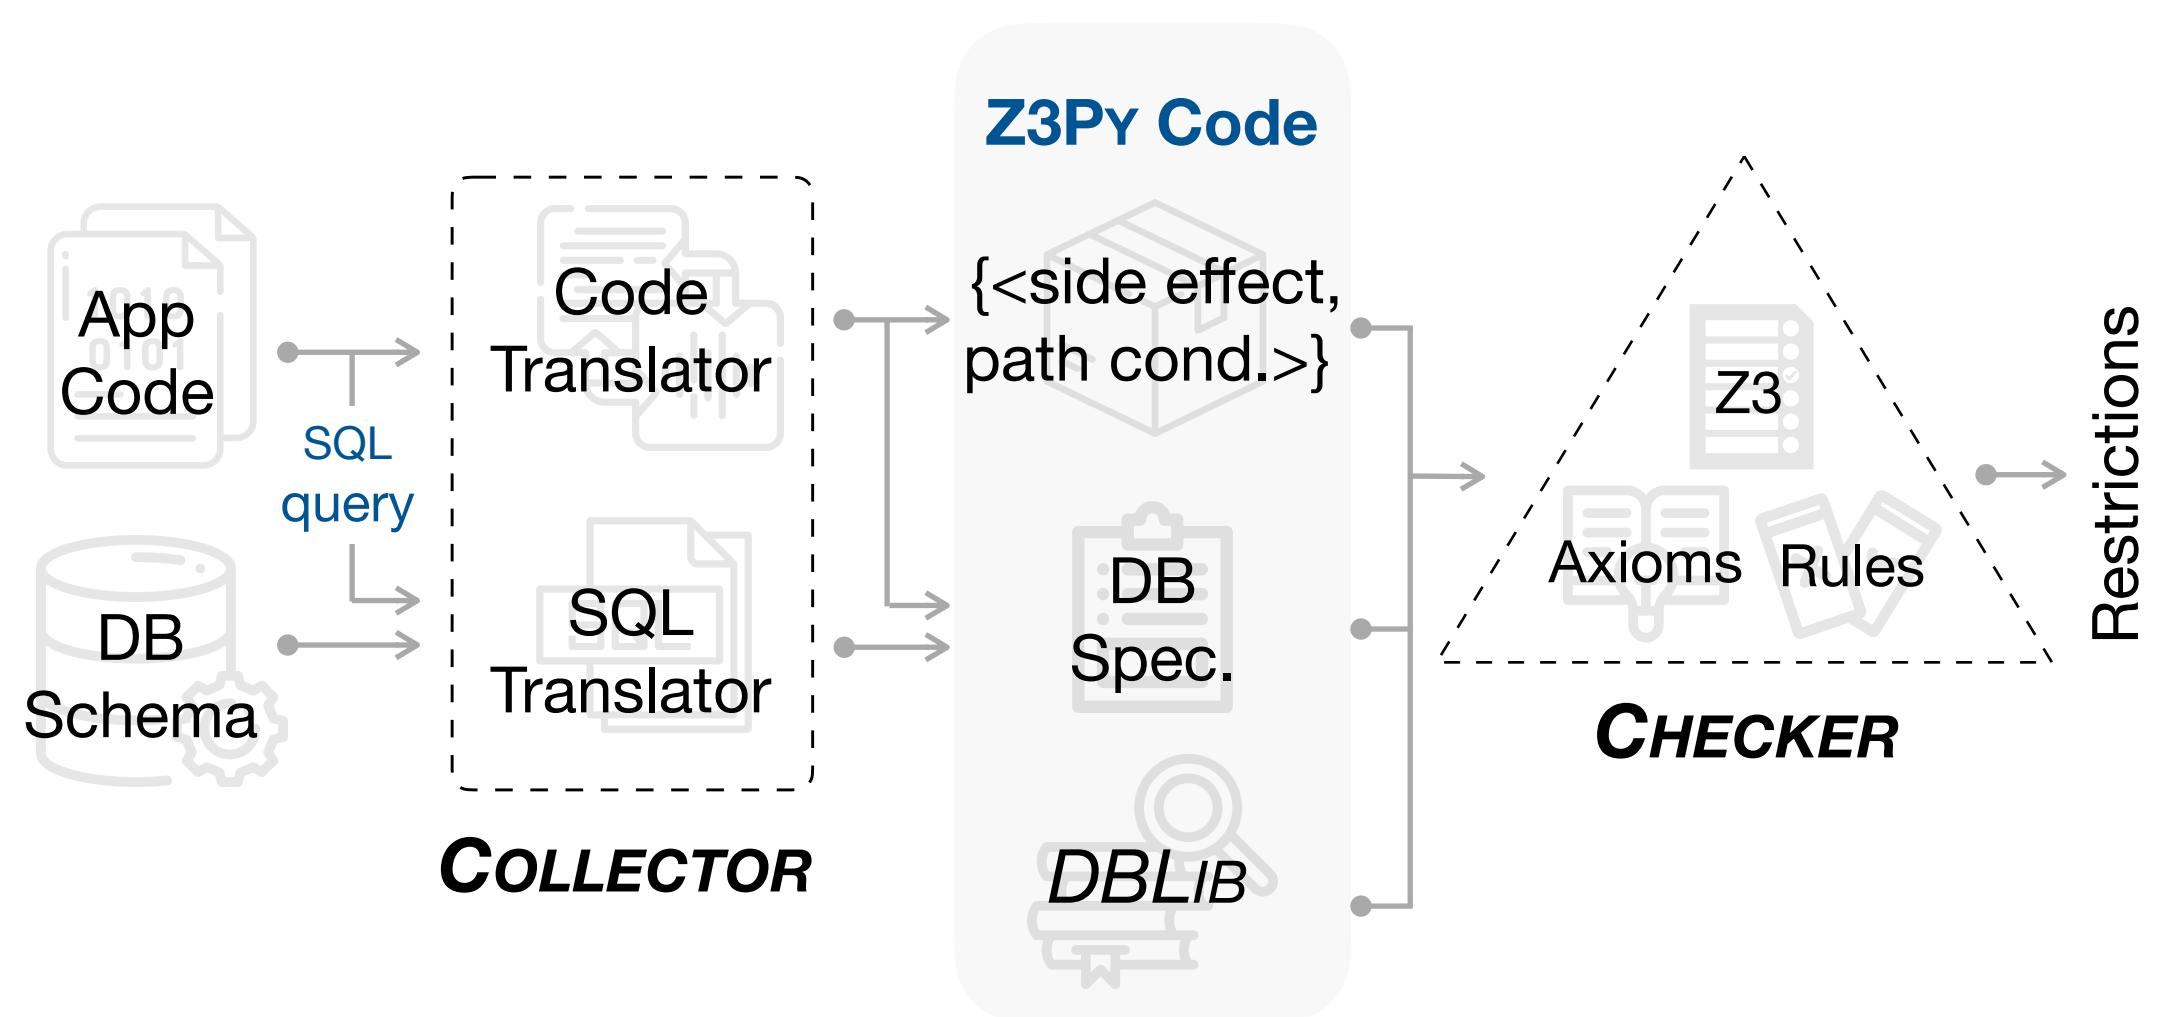
3-Site-AutoGR-Normal ✥

3-Site-Strong ■

Single-Site-Strong ━━━━ 3-Site-PoR ✕✕✕ 3-Site-AutoGR-Optimal ✕✕✕ 3-Site-AutoGR-Normal ━━ 3-Site-Strong ━＼＼







(Operation) $u \in \text{SideEffect} \times \text{PathCondition}$
(SideEffect) $\circ \in \text{State} \rightarrow \text{State}$
(PathCondition) $pcond \in (\text{State} \times \text{Argument}) \rightarrow \text{Bool}$
(Argument) $\text{arg} := \dots$
(State) $s := \dots$

